




جامعة الملك عبد الله  
للعلوم والتقنية  
King Abdullah University of  
Science and Technology




# Differentiable Rendering



Abdullah Hamdi , IVUL group meeting



# 3D representation



point cloud vs Voxels



**mesh**

# 3D representation



A Mesh is compact & useful for shading

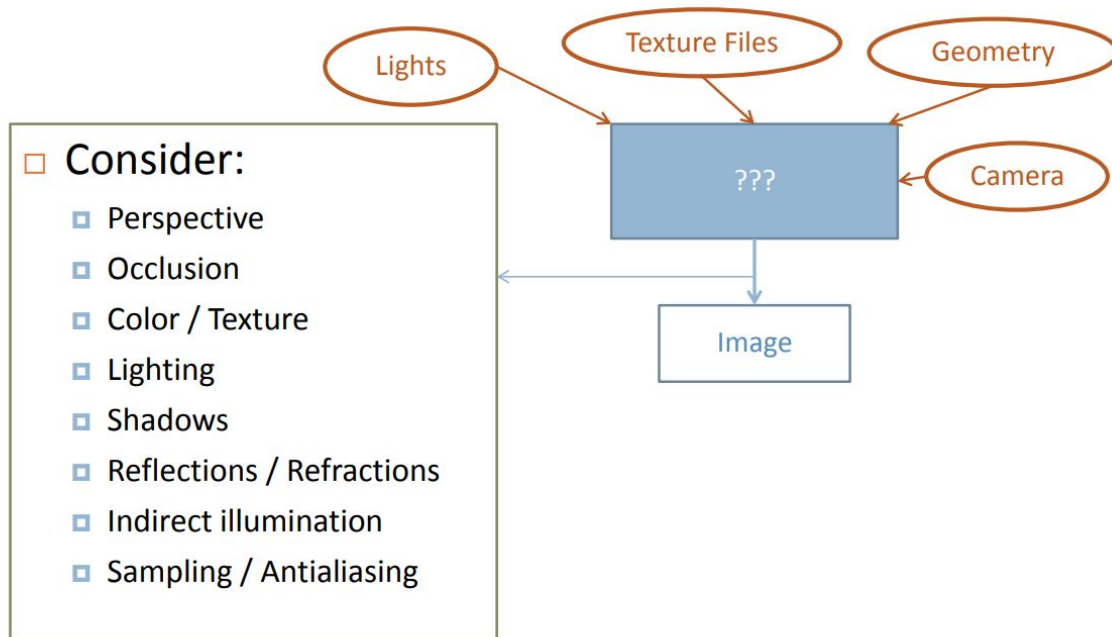


Yanir Kleiman , 2012

# Rendering in nutshell



## What is Rendering?



# Rendering in nutshell



## Two Approaches

- Start from **geometry**

- For each polygon / triangle:

- Is it visible?
    - Where is it?
    - What color is it?

**Rasterization**

- Start from **pixels**

- For each pixel in the final image:


- Which object is visible at this pixel?
    - What color is it?

**Ray Tracing**

# Shadows




## Rasterization vs ray tracing







# Degrees of realism



Polygonal model rendered in wire-frame  
(no visibility)



With visibility.



Shaded rendering. Note how the faces of the cube and cone have different intensities depending on their orientation relative to the light source.



Smooth patches and shading including highlights.



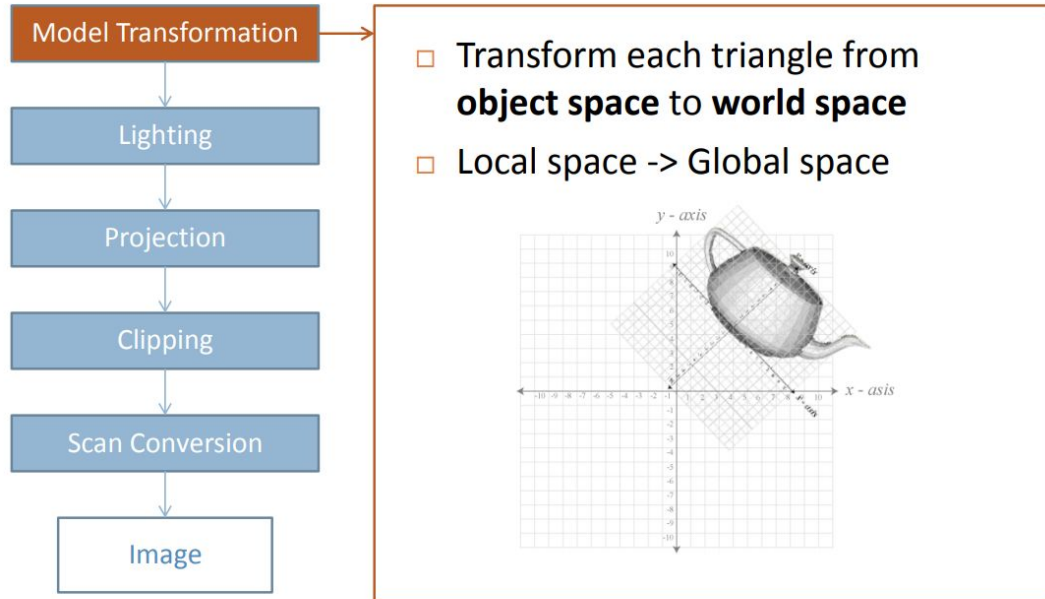
Texture-Mapping improves the appearance of surfaces (a better lighting is used too).



Shadows.



## Rasterization – Graphics Pipeline

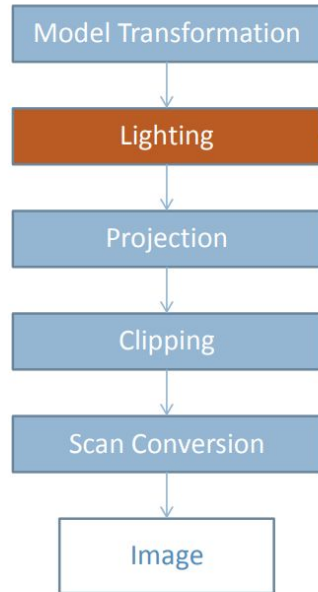





# Rendering in nutshell



## Rasterization – Graphics Pipeline



- Computation is based on angles between light source, object and camera (details later)

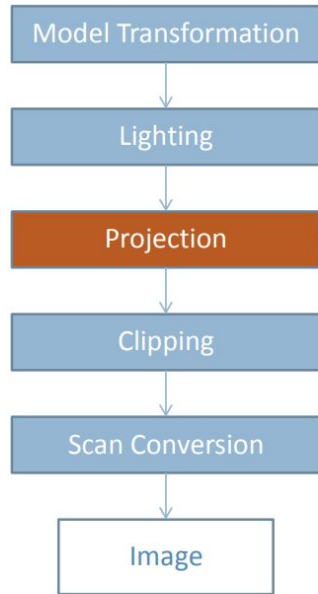


- Backface culling


# Rendering in nutshell



## Rasterization – Graphics Pipeline

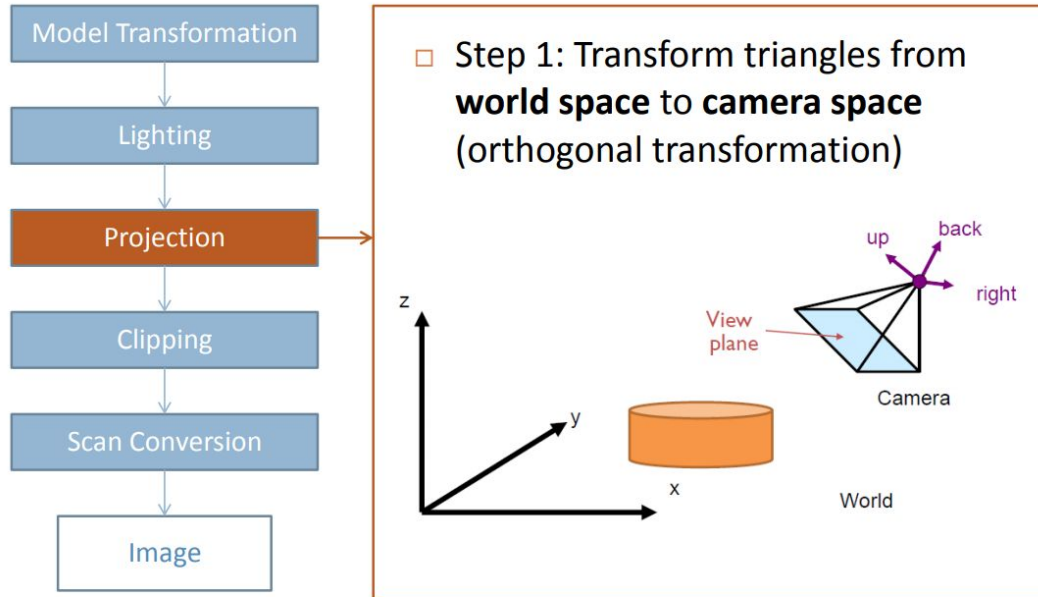


- Step 1: Transform triangles from **world space** to **camera space** (orthogonal transformation)



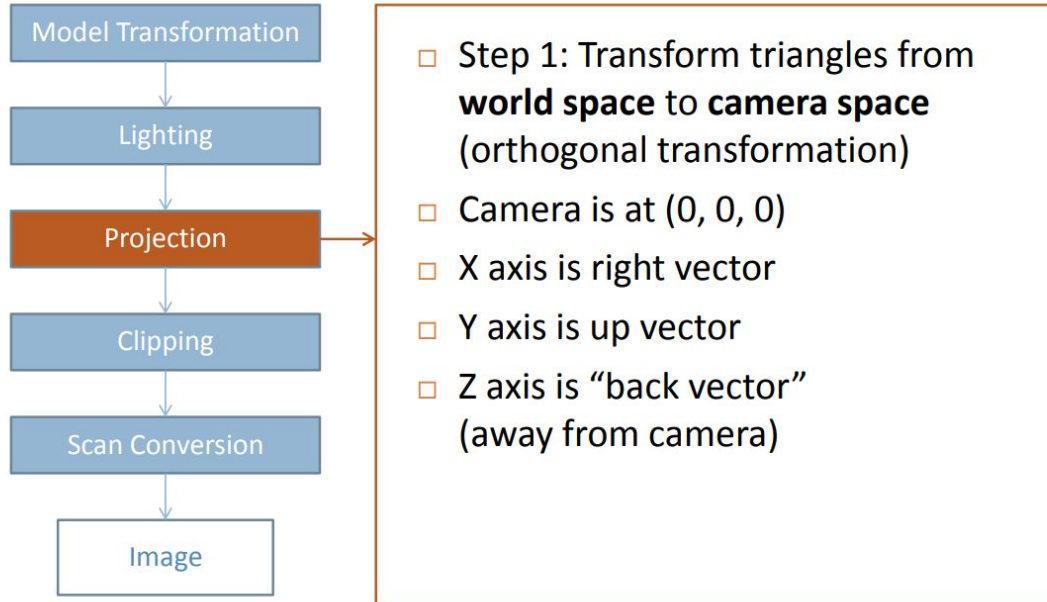


## Rasterization – Graphics Pipeline





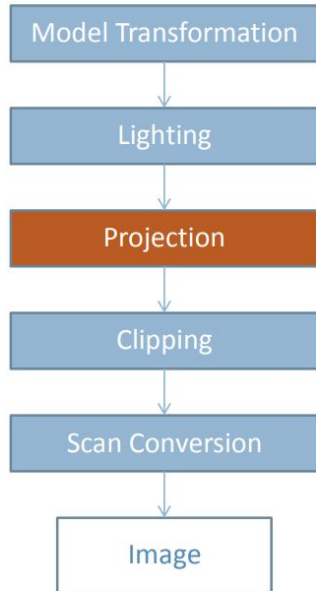
## Rasterization – Graphics Pipeline




# Rendering in nutshell



## Rasterization – Graphics Pipeline



- Step 2: Perspective Projection
- Depends on focal length ( $D$ )



- Calculate Z-Buffer

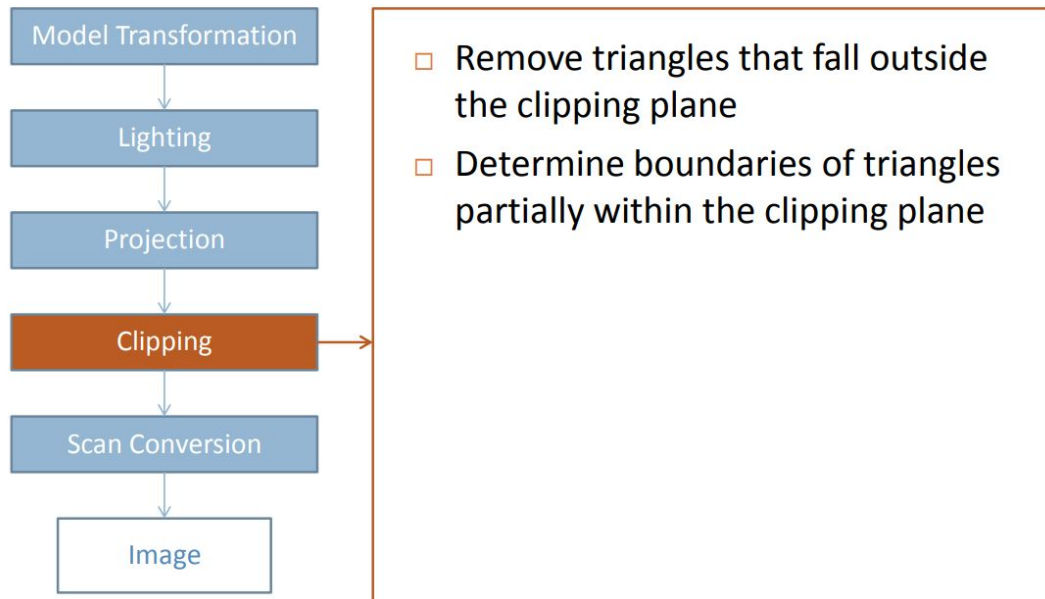


## Projection equation

$$\lambda \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \mathbf{P} \begin{bmatrix} u \\ v \\ w \\ 1 \end{bmatrix}, \quad \mathbf{P} = \mathbf{K} [\mathbf{R} \quad \mathbf{t}], \quad \mathbf{K} = \begin{bmatrix} f & 0 & c_x \\ 0 & f & c_y \\ 0 & 0 & 1 \end{bmatrix}$$



## Rasterization – Graphics Pipeline

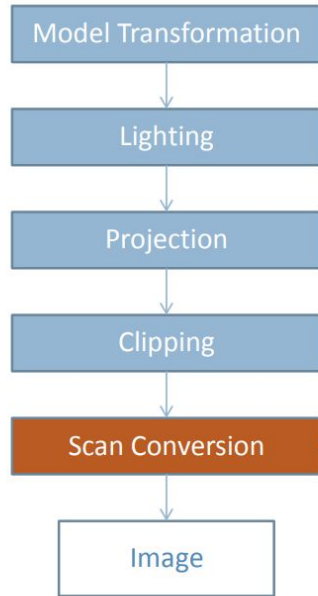





# Rendering in nutshell



## Rasterization – Graphics Pipeline



- Drawing the triangles in 2D
- Scanning horizontal scan lines for each triangle




# Rendering in nutshell



## Rasterization – Antialiasing

- Aliasing examples




# Rendering in nutshell



## Rasterization – Shadow Maps

- Render an image from the light's point of view (the light is the camera) ] **Shadow map**
- Keep “depth” from light of every pixel in the map
- During image render:  
Calculate position and depth on the **shadow map** for each **pixel** in the **final image** (not vertex!)
- If **pixel depth** > **shadow map depth** the pixel will not receive light from this source




# Rendering in nutshell



## Rasterization – Shadow Maps

- This solution is not optimal
  - Shadow map resolution is not correlated to render resolution – one shadow map pixel can span a lot of rendered pixels!
  - Shadow aliasing
  - Only allows sharp shadows
  - Semi-transparent objects
- Various hacks and complex solutions

Blurred hard shadows  
(shadow map)



True soft shadows  
(ray tracing)

# Neural 3D Mesh Renderer (CVPR'18)

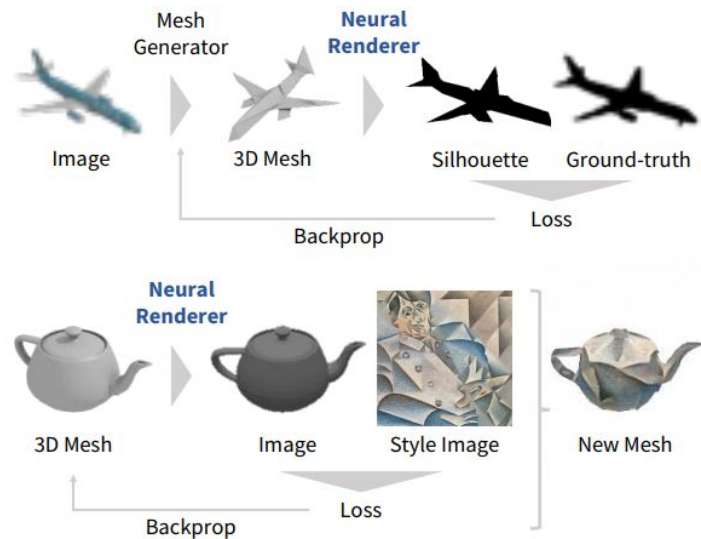



Figure 1. Pipelines for single-image 3D mesh reconstruction (upper) and 2D-to-3D style transfer (lower).



# Neural 3D Mesh Renderer (CVPR'18)




Figure 2. Illustration of our method.  $\mathbf{v}_i = \{x_i, y_i\}$  is one vertex of the face.  $I_j$  is the color of pixel  $P_j$ . The current position of  $x_i$  is  $x_0$ .  $x_1$  is the location of  $x_i$  where an edge of the face collides with the center of  $P_j$  when  $x_i$  moves to the right.  $I_j$  becomes  $I_{ij}$  when  $x_i = x_1$ .




Figure 3. Illustration of our method in the case where  $P_j$  is inside the face.  $I_j$  changes when  $x_i$  moves to the right or left.

# Neural 3D Mesh Renderer (CVPR'18)



$$\left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0} = \begin{cases} \frac{\delta_j^I}{\delta_i^x}; & \delta_j^P \delta_j^I < 0. \\ 0; & \delta_j^P \delta_j^I \geq 0. \end{cases}$$

$$\left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0} = \left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0}^a + \left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0}^b.$$

$$\left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0}^a = \begin{cases} \frac{\delta_j^{I^a}}{\delta_x^a}; & \delta_j^P \delta_j^{I^a} < 0. \\ 0; & \delta_j^P \delta_j^{I^a} \geq 0. \end{cases}$$

$$\left. \frac{\partial I_j(x_i)}{\partial x_i} \right|_{x_i=x_0}^b = \begin{cases} \frac{\delta_j^{I^b}}{\delta_x^b}; & \delta_j^P \delta_j^{I^b} < 0. \\ 0; & \delta_j^P \delta_j^{I^a} \geq 0. \end{cases}$$



# Neural 3D Mesh Renderer (CVPR'18)



$$I_j^l = (l^a + (\mathbf{n}^d \cdot \mathbf{n}_j) l^d) I_j.$$

$$\mathcal{L}_{\text{sl}}(x|\phi_i, s_i) = -\frac{|\hat{s}_i \odot s_i|_1}{|\hat{s}_i + s_i - \hat{s}_i \odot s_i|_1}.$$

$$\mathcal{L}_{\text{sm}}(x) = \sum_{\theta_i \in \mathcal{E}} (\cos \theta_i + 1)^2.$$

# Neural 3D Mesh Renderer (CVPR'18)




	airplane	bench	dresser	car	chair	display	lamp
Retrieval [36]	0.5564	0.4875	0.5713	0.6519	0.3512	0.3958	0.2905
Voxel-based [36]	0.5556	0.4924	0.6823	<b>0.7123</b>	0.4494	0.5395	<b>0.4223</b>
Mesh-based (ours)	<b>0.6172</b>	<b>0.4998</b>	<b>0.7143</b>	0.7095	<b>0.4990</b>	<b>0.5831</b>	0.4126

	loudspeaker	rifle	sofa	table	telephone	vessel	mean
Retrieval [36]	0.4600	0.5133	0.5314	0.3097	0.6696	0.4078	0.4766
Voxel-based [36]	0.5868	0.5987	0.6221	<b>0.4938</b>	0.7504	0.5507	0.5736
Mesh-based (ours)	<b>0.6536</b>	<b>0.6322</b>	<b>0.6735</b>	0.4829	<b>0.7777</b>	<b>0.5645</b>	<b>0.6016</b>

Table 1. Reconstruction accuracy measured by voxel IoU. Higher is better. Our mesh-based approach outperforms the voxel-based approach [36] in 10 out of 13 categories.

# Neural 3D Mesh Renderer (CVPR'18)



applications






Deep dream



applications




# Style transfer






## Optimizing mesh to match a projection






## Optimizing texture







## Optimizing view point



# Jupyter Notebook !!



## Let's all do it !!!

```
Neural_renderer
Python [conda env:mytorch]

Neural Mesh Renderer in pytorch
Kato etal (CVPR18)
implements four examples to use Neaural Renderer with pytorch

-----
go to the repo here and hit
sudo python setup.py install

importing relevent modules
In [1]: from __future__ import division
import os
import argparse
import glob

import torch
import torch.nn as nn
import numpy as np
from skimage.io import imread, imsave
from tqdm import tqdm_notebook as tqdm

import imageio
from IPython.display import display, HTML

import neural_renderer as nr

current_dir = os.getcwd()
data_dir = os.path.join(current_dir, 'data/3d_renderer')

supporting functions
In [2]: def make_gif(filename):
with imageio.get_writer(filename, mode='I') as writer:
for filename in sorted(glob.glob('/tmp/tmp_*.png')):
writer.append_data(imread(filename))
```

