

Imitation Learning (Behavior Cloning)

Fabian Caba Heilbron
IVUL – February 16

http://rll.berkeley.edu/deeprlcourse/docs/week_2_lecture_1_behavior_cloning.pdf



Imitation vs Reinforcement



Define a reward for each task!

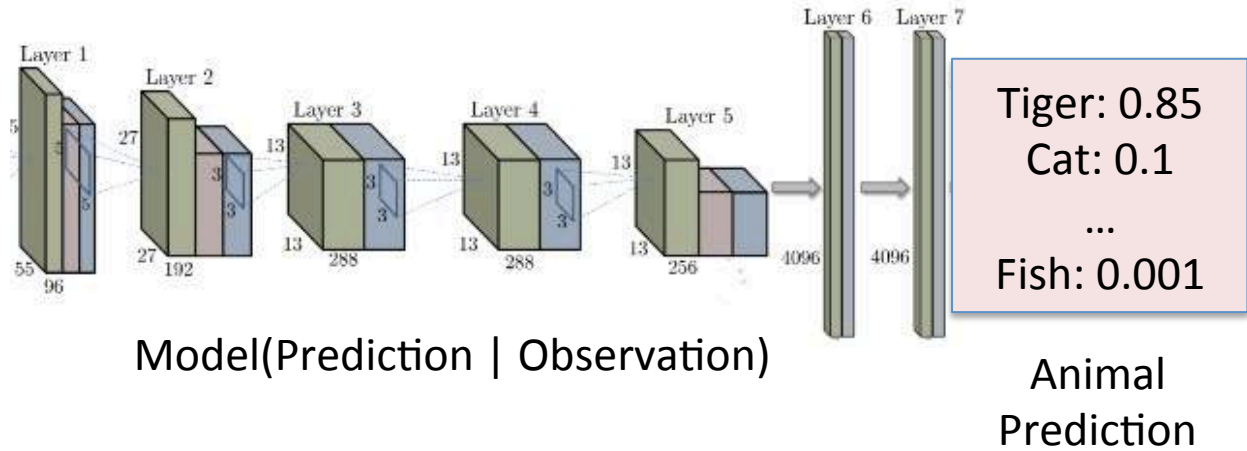
Today's talk

- Definition of sequential problem
- Imitation learning: supervised learning for decision making
 - Does imitation learning works?
 - How can we make it work more often?
- Case studies of recent work in (deep) imitation learning

Supervised Learning




Observation



Supervised Learning



Observation




Model(Prediction | Observation)

Scene
Prediction

Supervised Learning



Observation



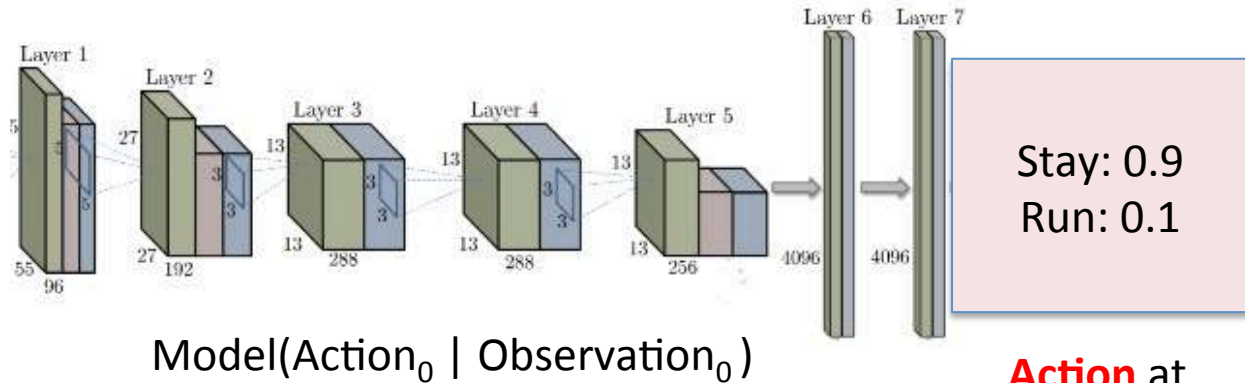
Model(Prediction | Observation)

Safety
Prediction

Sequential decision problem



Observation at
time 0

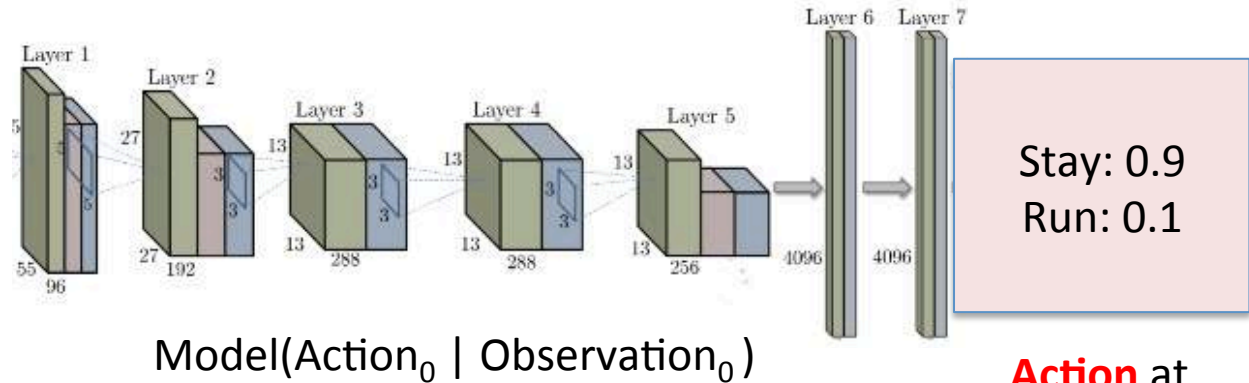


Action at
time 0

Sequential decision problem



Observation at
time 0



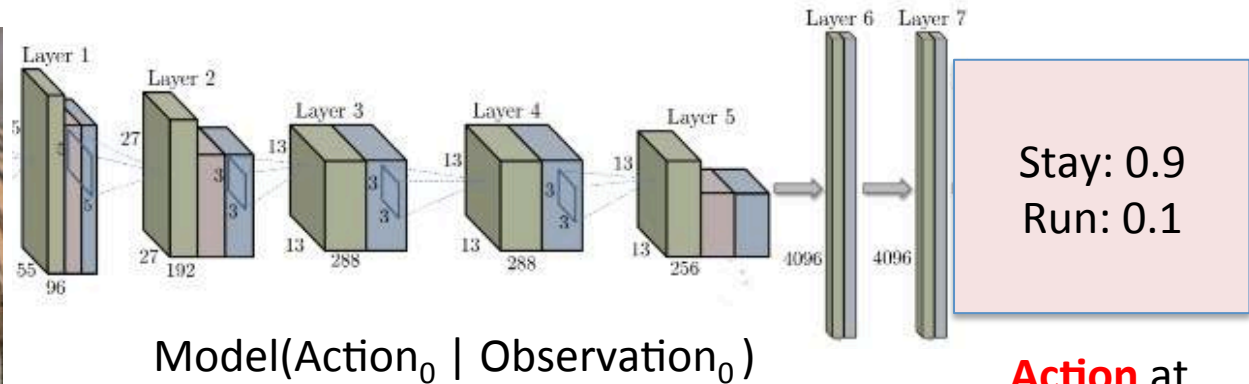
Action at
time 0



Sequential decision problem



Observation at
Time 1




Action at
Time 0



Sequential decision problem



Observation at
Time 1




Action at
Time 1

Sequential decision problem



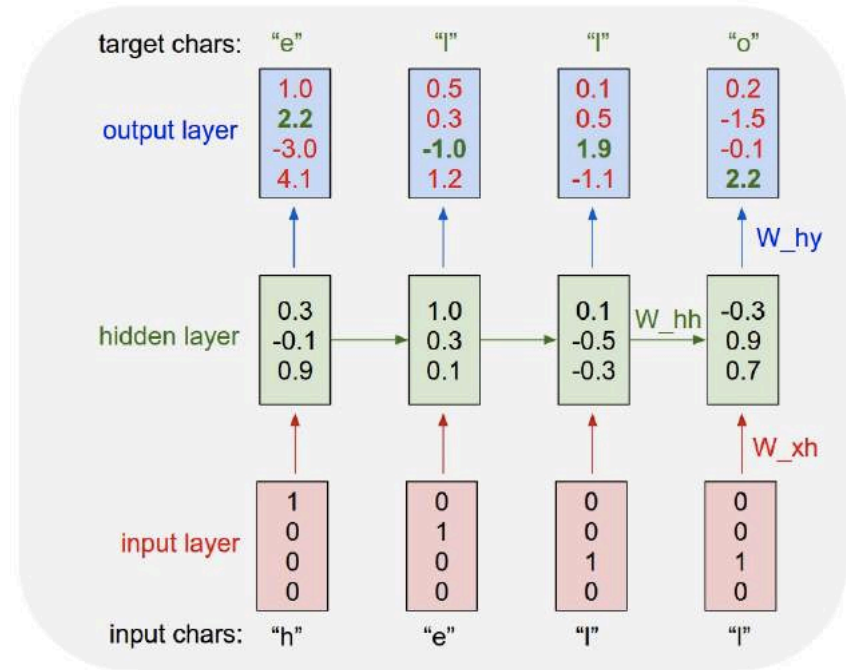
Observation at
Time t



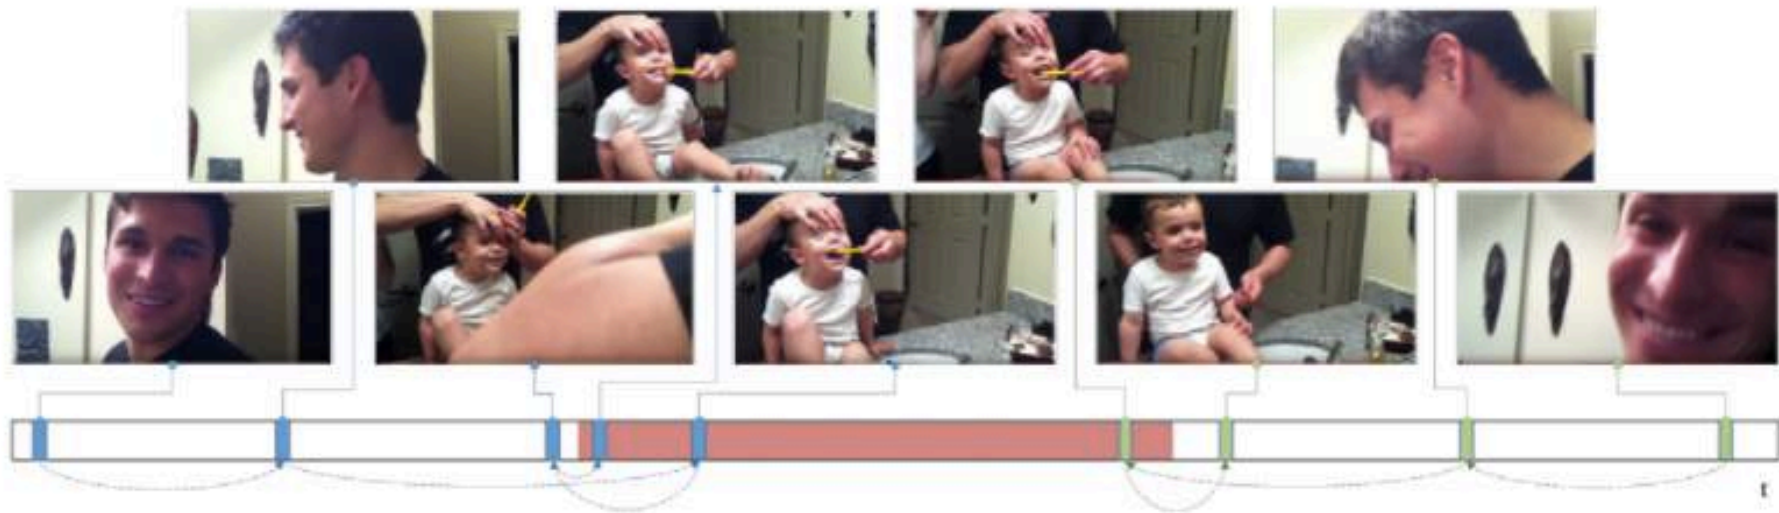
$\text{Model}(\text{Action}_t \mid (\text{Observation}_t, \text{Action}_{t-1}))$

Action at
Time t

Sequential decision problem examples



Sequential decision problem examples




Learning to search actions


Today's talk

- ~~Definition of sequential problem~~
- Imitation learning: supervised learning for decision making
 - Does imitation learning works?
 - How can we make it work more often?
- Case studies of recent work in (deep) imitation learning

Steering wheel prediction



Steering wheel prediction




Collect observation and actions

Steering wheel prediction



Observation at time t



Action at time t




Collect observation and actions




Should it work?

Should it work? **NO**

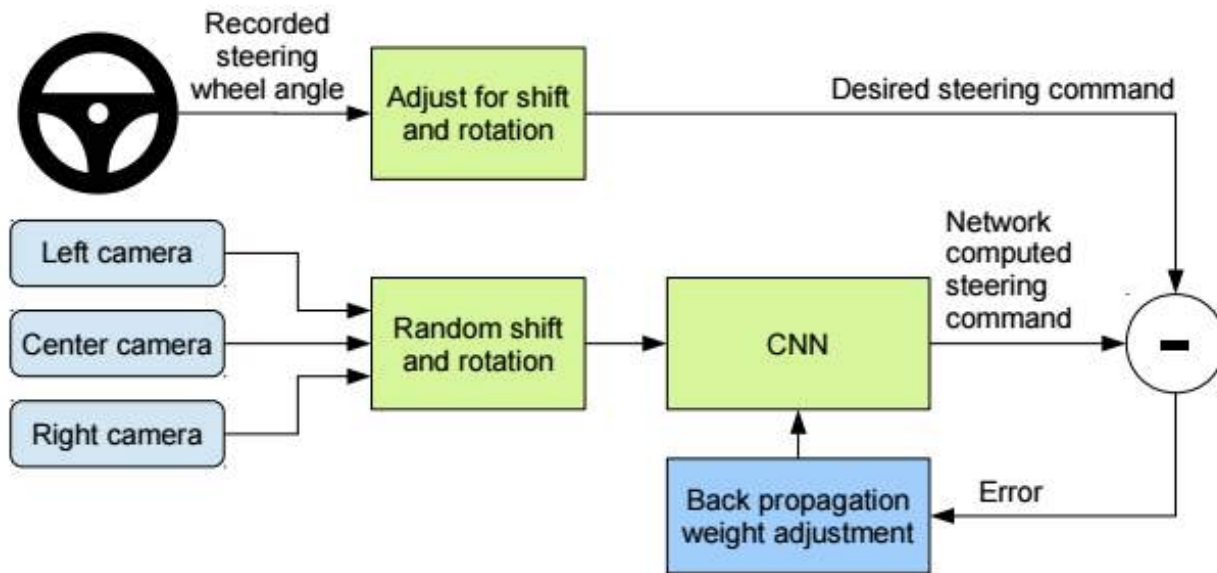


Should it work? **NO**

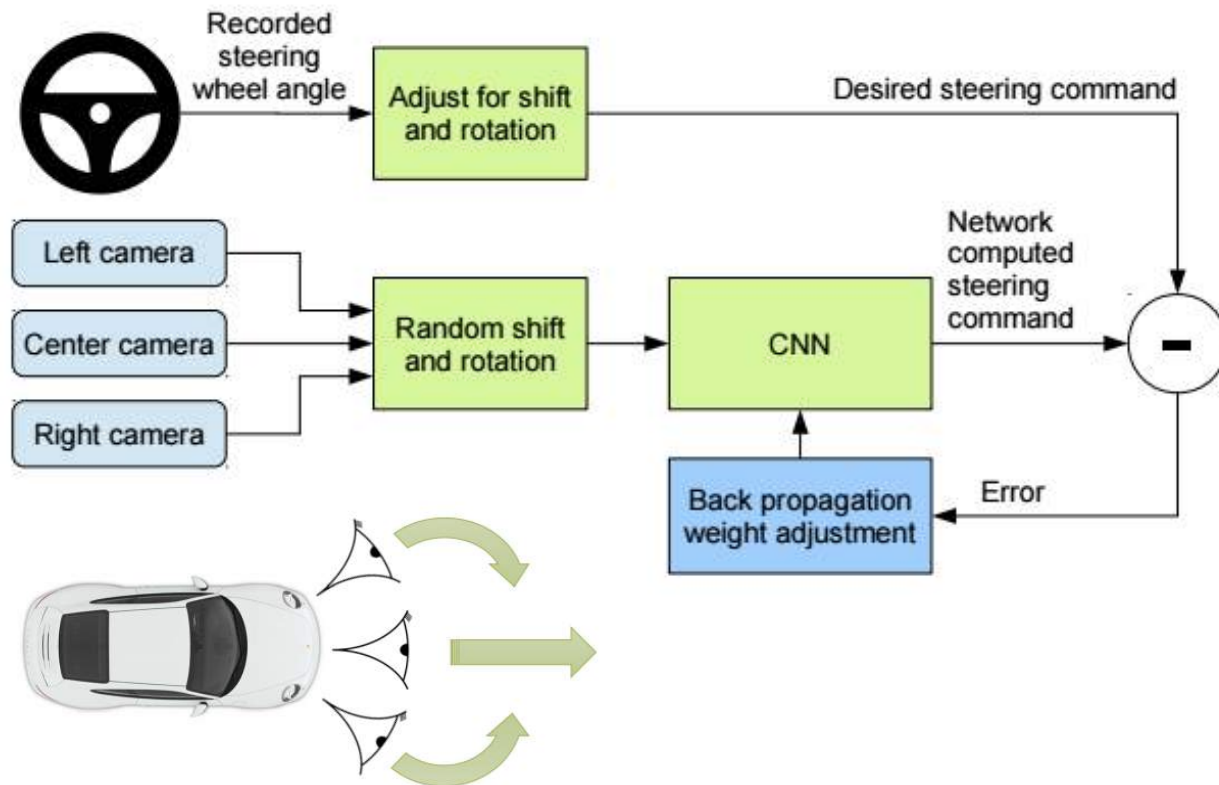


Does it work? **Yes!** NVidia

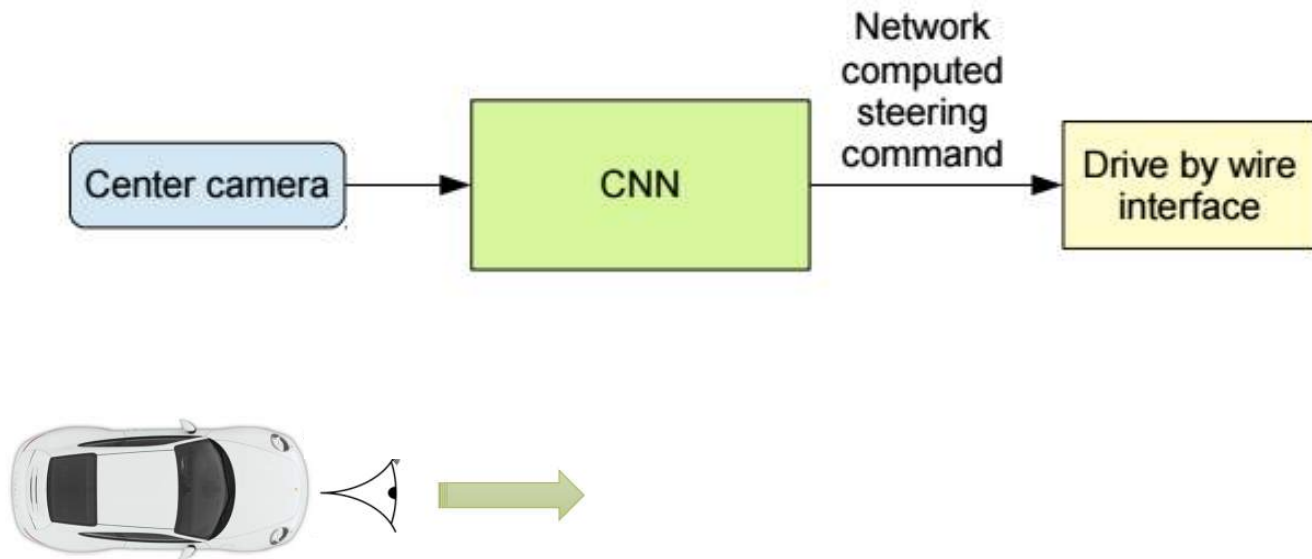
Why it works?




Why it works? It's a Hack!



Why it works? It's a Hack!



Can we make it work more often?




Can we make it work more often?

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

- 
1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
 2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_π with actions \mathbf{u}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Dagger: Watch!


Can we make it work more often?

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$


how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

- 
1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
 2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_π with actions \mathbf{u}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$



Data as Demonstrator




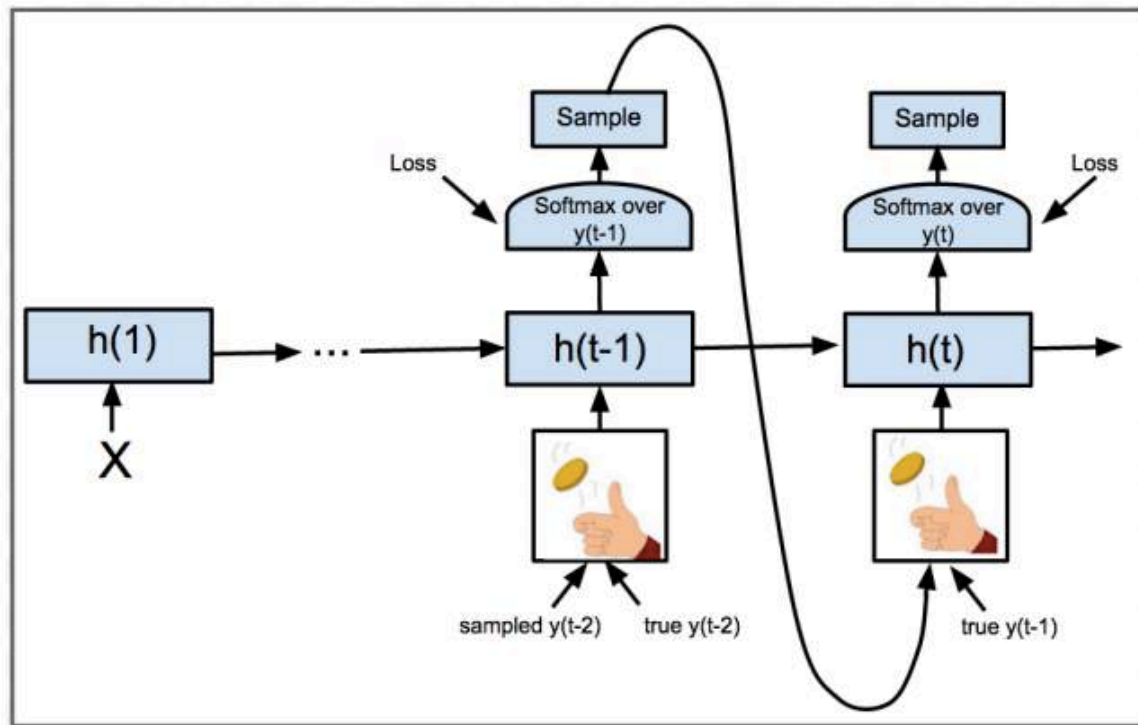
Today's talk

- ~~Definition of sequential problem~~
- ~~Imitation learning: supervised learning for decision making~~
 - ~~Does imitation learning works?~~
 - ~~How can we make it work more often?~~
- Case studies of recent work in (deep) imitation learning

Case Study 1: Trail following as classification

- Watch!

Case Study 2: Scheduled Sampling



Imitation Learning Summary

- Some times works well
 - Hacks
 - Intelligent way of getting training samples
 - Really diverse dataset that summarizes the data distribution
- Problems
 - Humans needs to provide data, which is usually finite
 - Humans are not good at provided some kind of actions

That's all Folks!