Siamese Network: Architecture and Applications in Computer Vision

February 9, 2017

Outline

- Metric Learning
- Siamese Architecture
- Siamese Network: Applications in computer vision

Traditional Classification

- All classes should be known
- Number of samples per category should be large
- Training samples should have a known class

Distance-based Methods

- Compute a similarity metric between the pattern to be classified or verified and a library of stored prototypes

- Can be used to compare or match new samples from previously-unseen categories (e.g. faces from people not seen during training)

Metric Learning

- Euclidean distance *vs* Mahalanobis distance

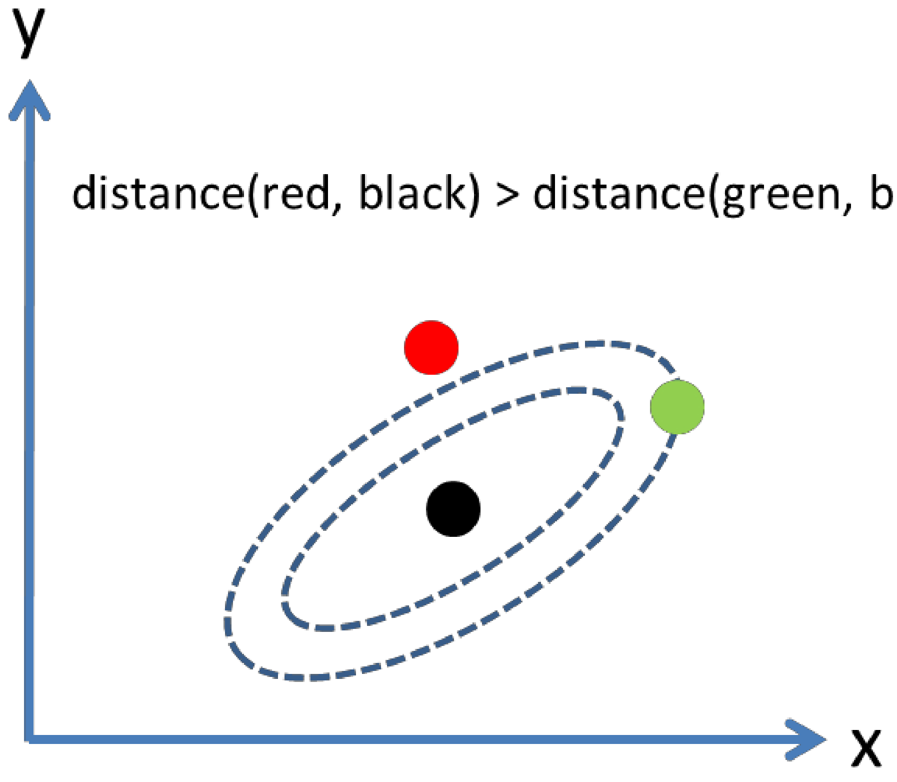

http://horicky.blogspot.jp/2012/08/measuring-similarity-and-distance.html

Metric Learning

Mahalanobis Distance Metric Learning

- ***Euclidean distance***
- ***Mahalanobis distance*** $d(\vec{x},\vec{y}) = \sqrt{(\vec{x}-\vec{y})^T S^{-1} (\vec{x}-\vec{y})}.$
- ***Mahalanobis Distance Metric Learning***

$$d(x,y) = d_A(x,y) = ||x-y||_A = \sqrt{(x-y)^T A (x-y)}$$

$$\begin{array}{ll} \min\limits_{A} & \sum_{(x_i,x_j)\in\mathcal{S}} ||x_i - x_j||_A^2 \\ \text{s.t.} & \sum_{(x_i,x_j)\in\mathcal{D}} ||x_i - x_j||_A \geq 1 \\ & A \succeq 0 \end{array}$$

Xing E P,Jordan M I, Russell S, et al. Distance metric learning with application to clustering with side-information[C], NIPS2002: 505-512.

Metric Learning
Large-Margin Nearest Neighbors(LMNN)

$$\min_{A \succeq 0} \sum_{(i,j)\in\mathcal{S}} d_A(\boldsymbol{x}_i, \boldsymbol{x}_j) + \lambda \sum_{(i,j,k)\in\mathcal{R}} [1 + d_A(\boldsymbol{x}_i, \boldsymbol{x}_j) - d_A(\boldsymbol{x}_i, \boldsymbol{x}_k)]_+$$

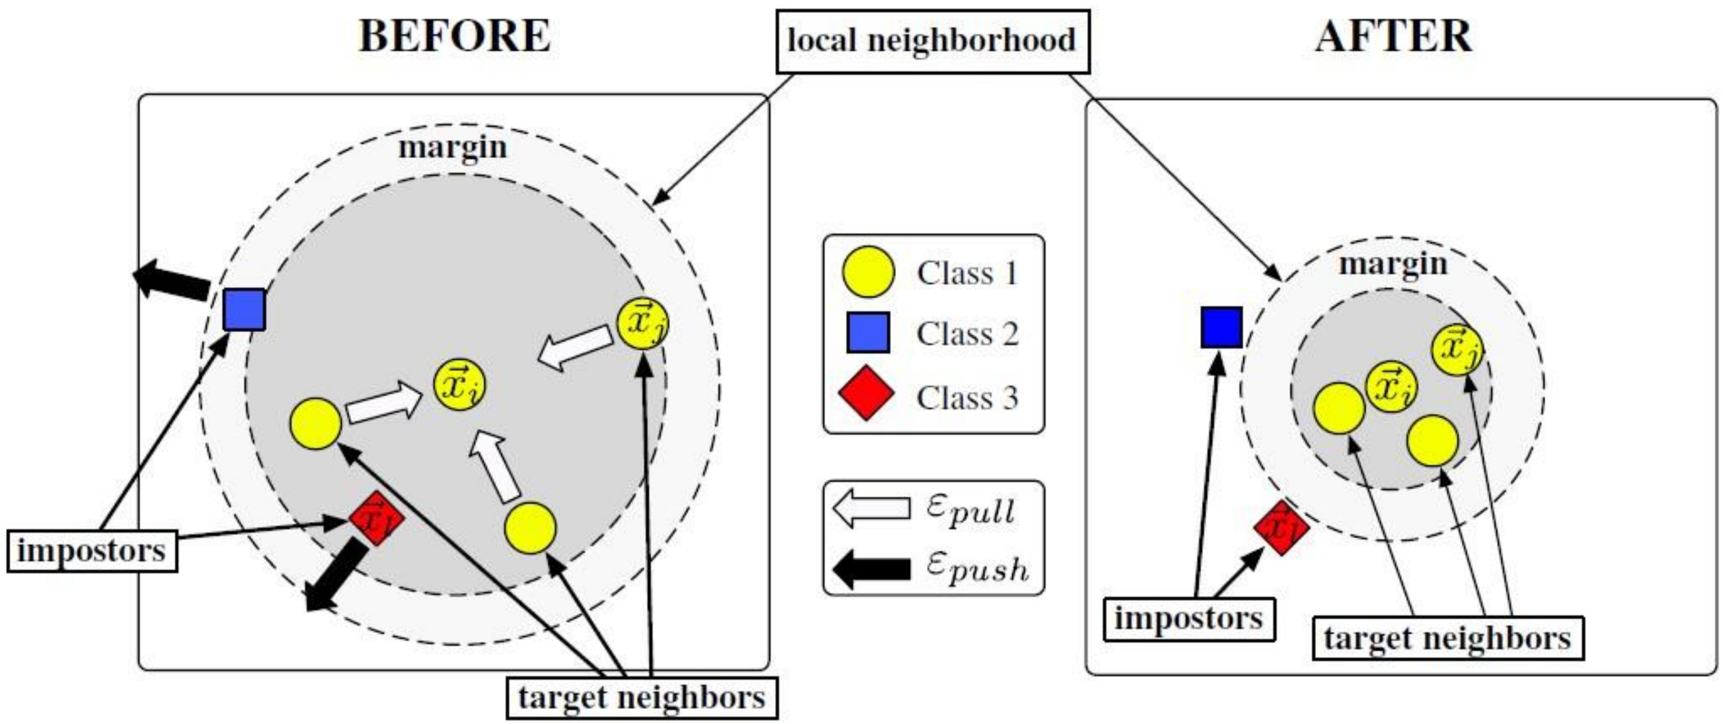

Brian Kulis, Metric Learning: A Survey. *web.cse.ohio-state.edu/~kulis/pubs/ftml_**metric_learning**.pdf*

Siamese

- Someone or something from Thailand:
  - The Thai language, The Thai people
- Siamese, an informal term for conjoined or fused:
  - Siamese twins, conjoined twins
  - Siamesing (engineering), the practice, whose name is derived from siamese twins, of combining two devices (such as cylinder ports or cooling jackets) together into a closely coupled pair, so as to save space between them.

http://en.wikipedia.org/wiki/Siamese

Siamese Architecture

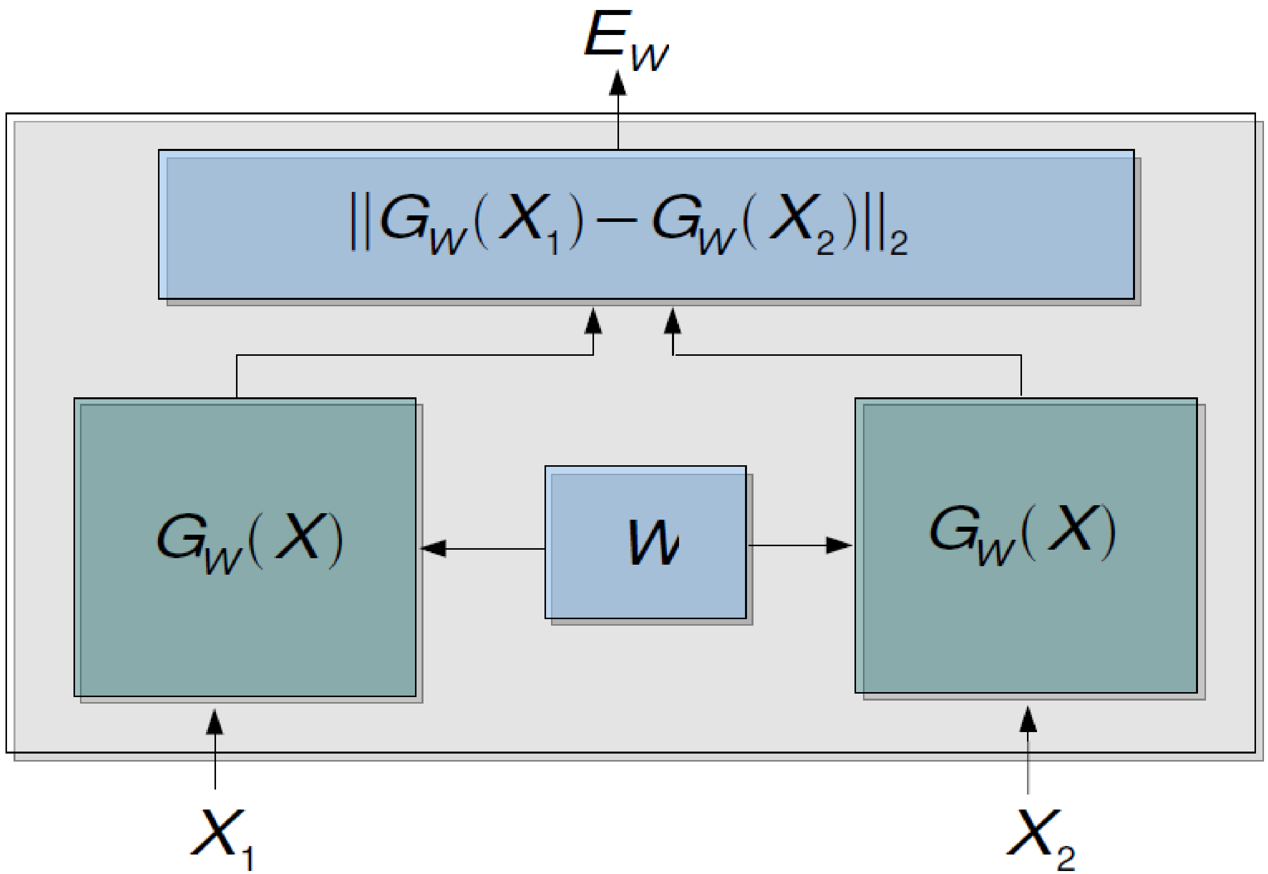

Learning Hierarchies of *Invariant Features*. Yann LeCun.
helper.ipam.ucla.edu/publications/gss2012/gss2012_10739.pdf

Siamese Architecture and loss function

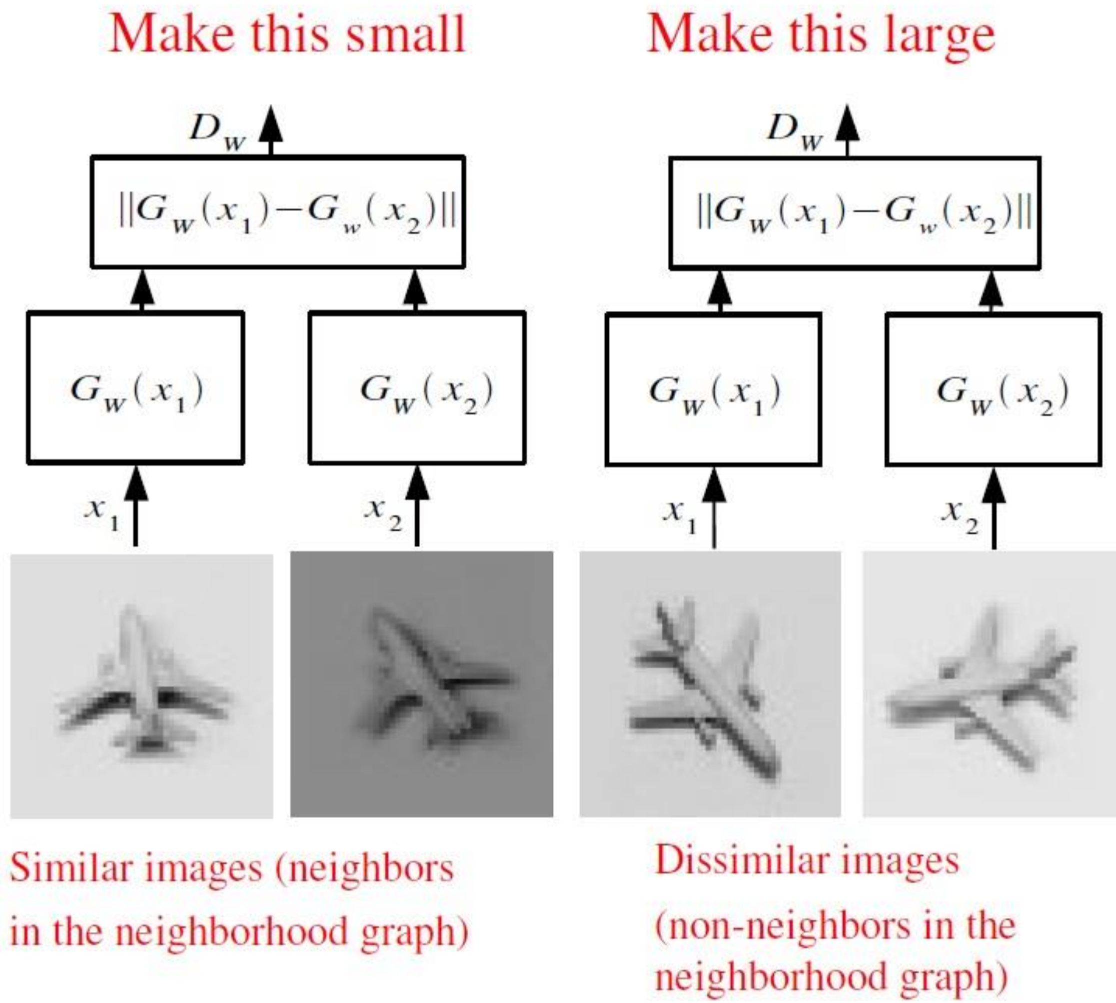

Learning Hierarchies of *Invariant Features*. Yann LeCun.
helper.ipam.ucla.edu/publications/gss2012/gss2012_10739.pdf

Loss function

$$L_{similar} = \frac{1}{2} D_W^2 \qquad L_{dissimilar} = \frac{1}{2} \{max(0, m - D_W)\}^2$$

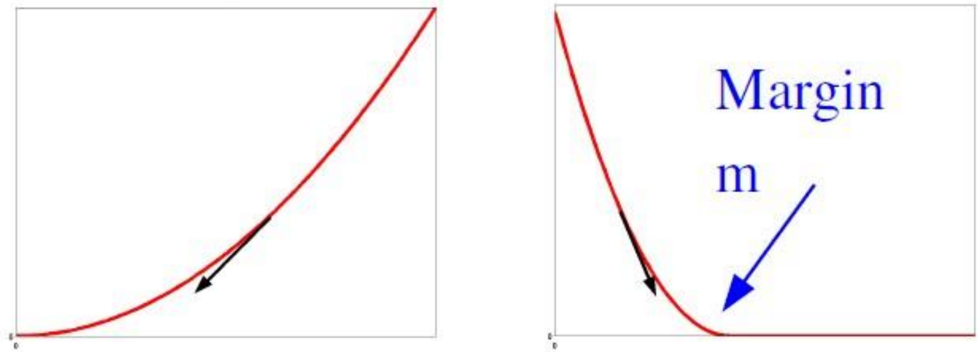

Hinge Loss

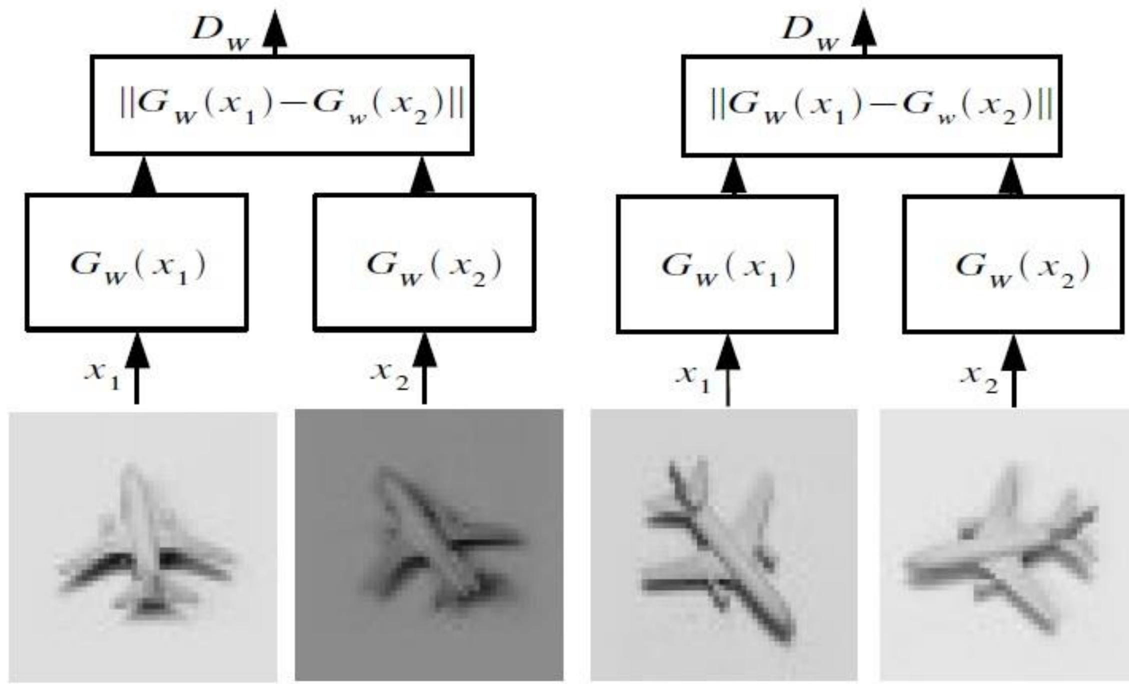

Learning Hierarchies of *Invariant Features*. Yann LeCun.
helper.ipam.ucla.edu/publications/gss2012/gss2012_10739.pdf

Siamese Network

Application in Signature Verification

- The input is 8(feature) x 200(time) units.
- The cosine distance was used, (1 for genuine pairs, -1 for forgery pairs)

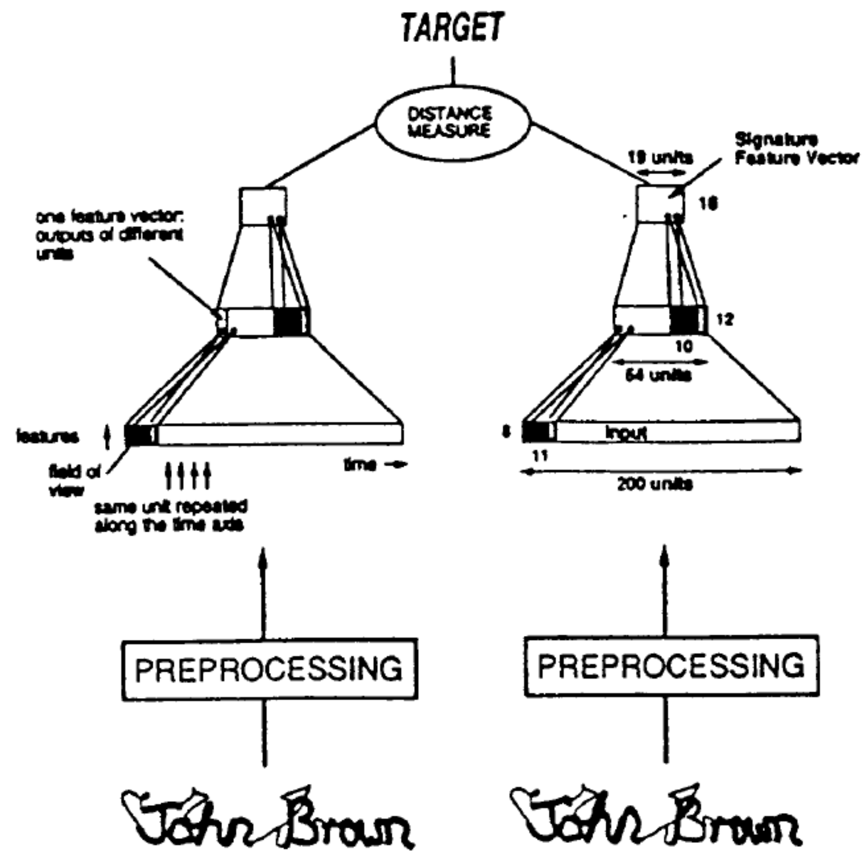

Bromley J, Guyon I, Lecun Y, et al. Signature Verification using a" Siamese" Time Delay Neural Network, NIPS Proc. 1994.

Siamese Network
Application in Dimensionality reduction

The exact loss function is

$$L(W, Y, \vec{X}_1, \vec{X}_2) = (1 - Y)\frac{1}{2}(D_W)^2 + (Y)\frac{1}{2}\{max(0, m - D_W)\}^2$$

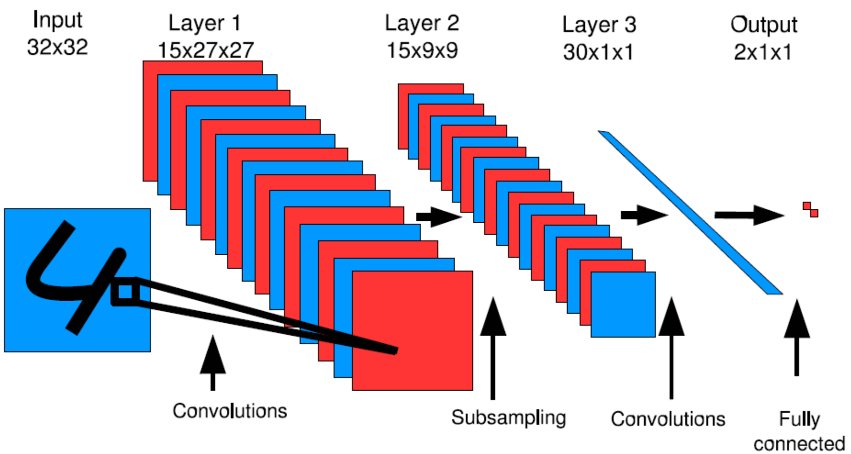

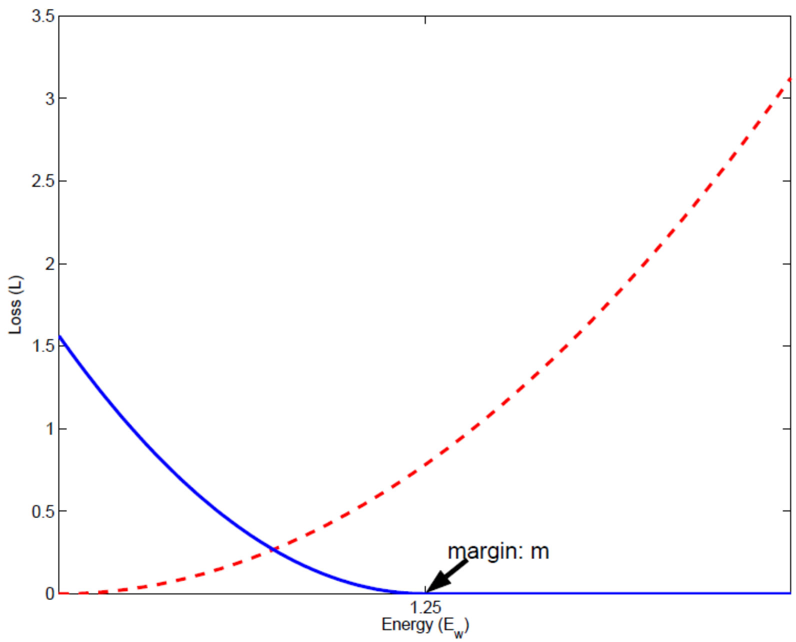

Hadsell R, Chopra S, LeCun Y. Dimensionality reduction by learning an invariant mapping, CVPR 2006

Siamese Network

Application in Dimensionality reduction

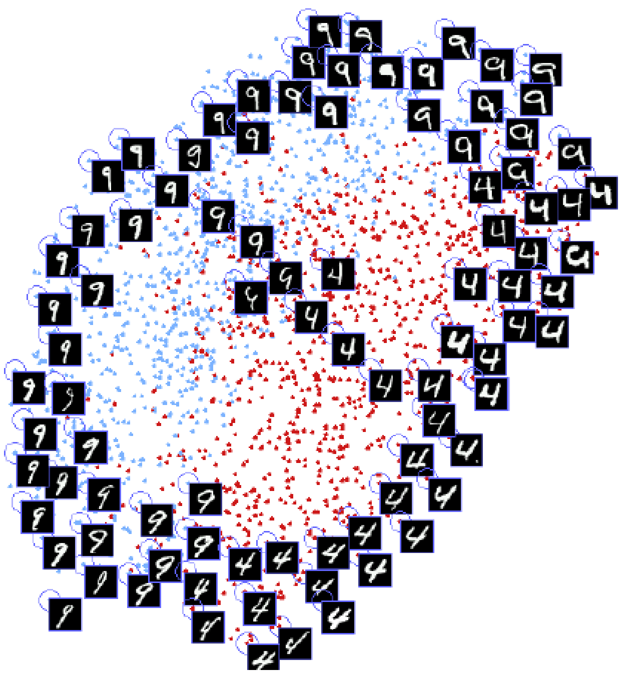

LearnedMapping of MNIST samples

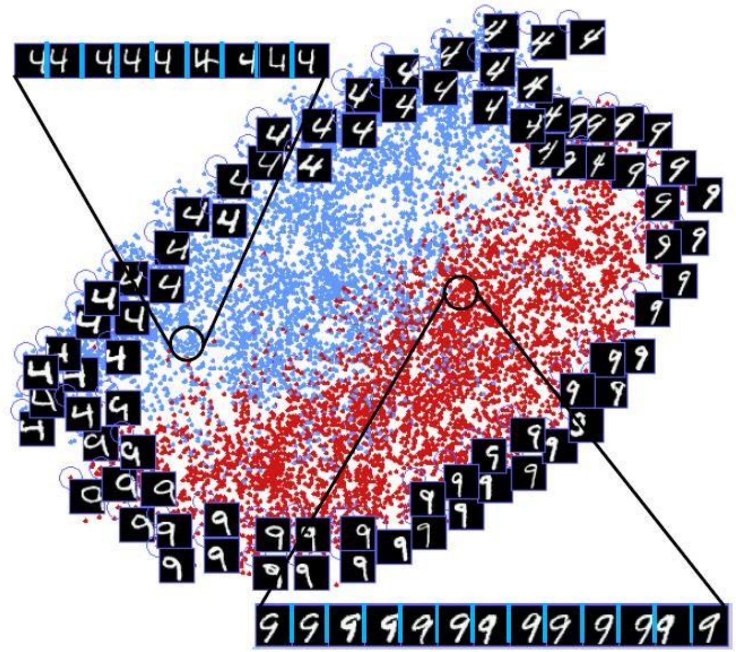

Learning a Shift Invariant Mapping of MNIST samples

Hadsell R, Chopra S, LeCun Y. Dimensionality reduction by learning an invariant mapping, CVPR 2006

Image Descriptors

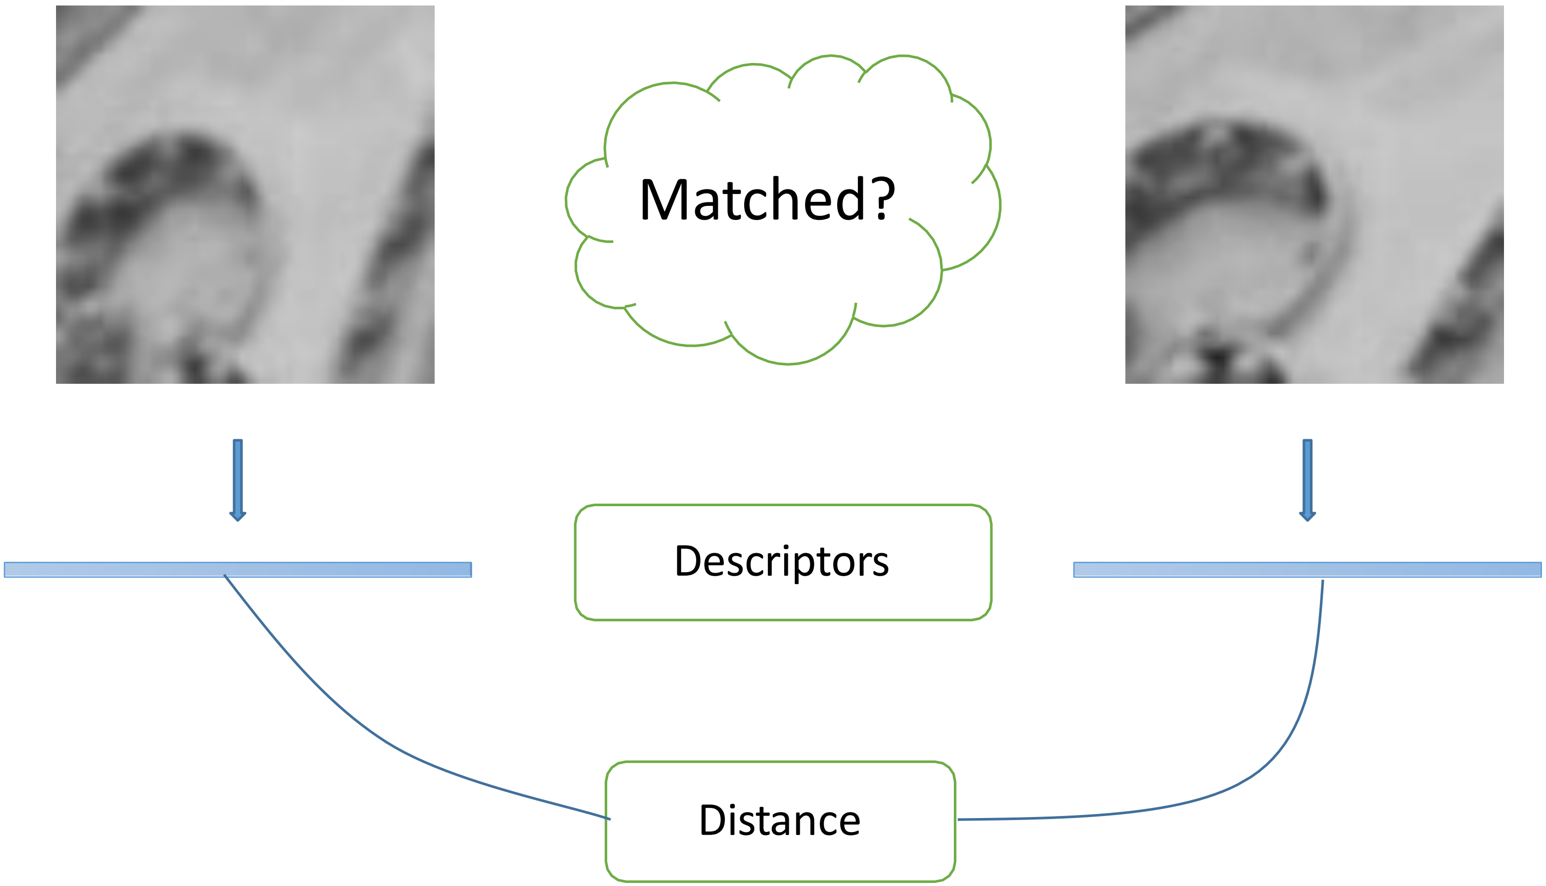

Siamese Network

Application in Learning Image Descriptors (I)

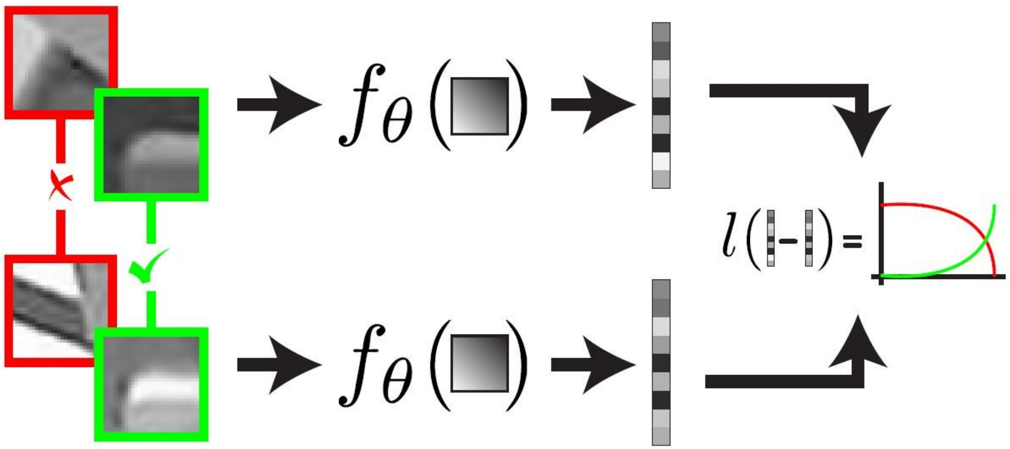

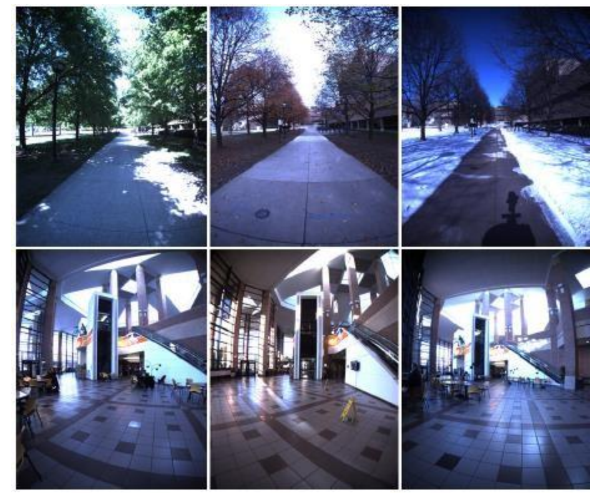

Using the contrastive cost function

$$l_{\boldsymbol{\theta}}\left(\mathbf{y}_i, \mathbf{y}_j\right) = \begin{cases} s_{ij} d_{ij}^2, & \text{if matching} \\ \max\left(1.0 - d_{ij}^2,\ 0\right), & \text{if non-matching} \end{cases}$$

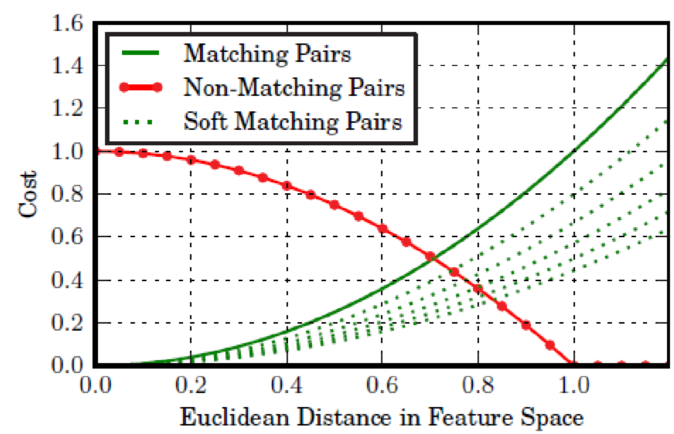

Nicholas Carlevaris-Bianco and Ryan M. Eustice, Learning visual feature descriptors for dynamic lighting conditions. In Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems, 2014

Siamese Network

Application in Learning Image Descriptors (I)

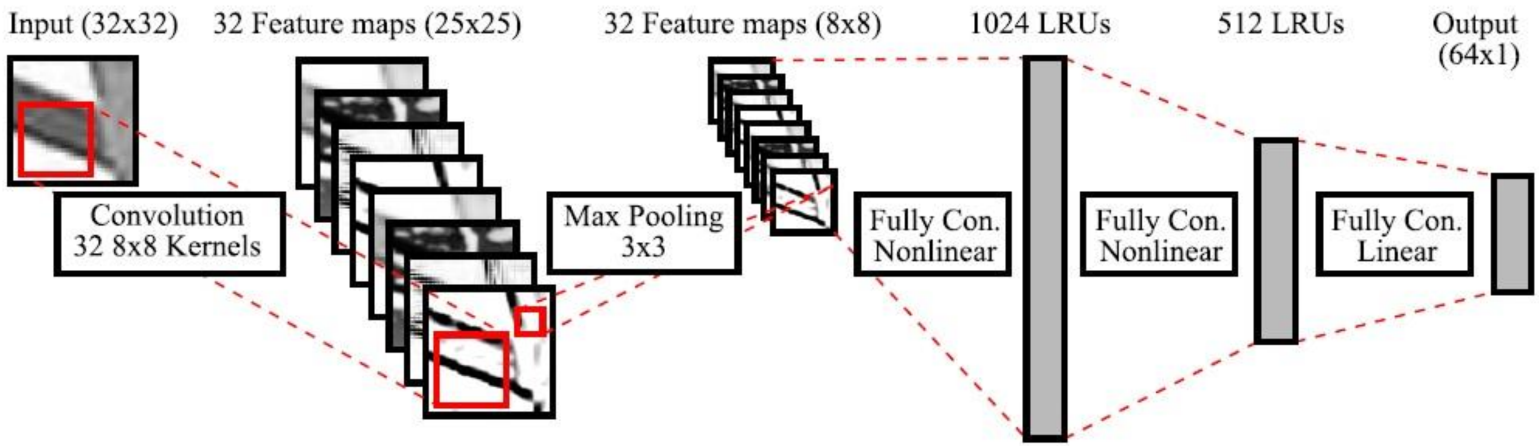

CNN Model

Nicholas Carlevaris-Bianco and Ryan M. Eustice, Learning visual feature descriptors for dynamic lighting conditions. In Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems, 2014

Siamese Network

Application in Learning Image Descriptors (II)

$$d_D(\mathbf{x}_1, \mathbf{x}_2) = \|D(\mathbf{x}_1) - D(\mathbf{x}_2)\|_2$$

$$l(\mathbf{x}_1, \mathbf{x}_2, \delta) = \delta \cdot l_P(d_D(\mathbf{x}_1, \mathbf{x}_2)) + (1 - \delta) \cdot l_N(d_D(\mathbf{x}_1, \mathbf{x}_2))$$

$$l_P(d_D(\mathbf{x}_1, \mathbf{x}_2)) = d_D(\mathbf{x}_1, \mathbf{x}_2)$$

$$l_N(d_D(\mathbf{x}_1, \mathbf{x}_2)) = \max(0, m - d_D(\mathbf{x}_1, \mathbf{x}_2))$$

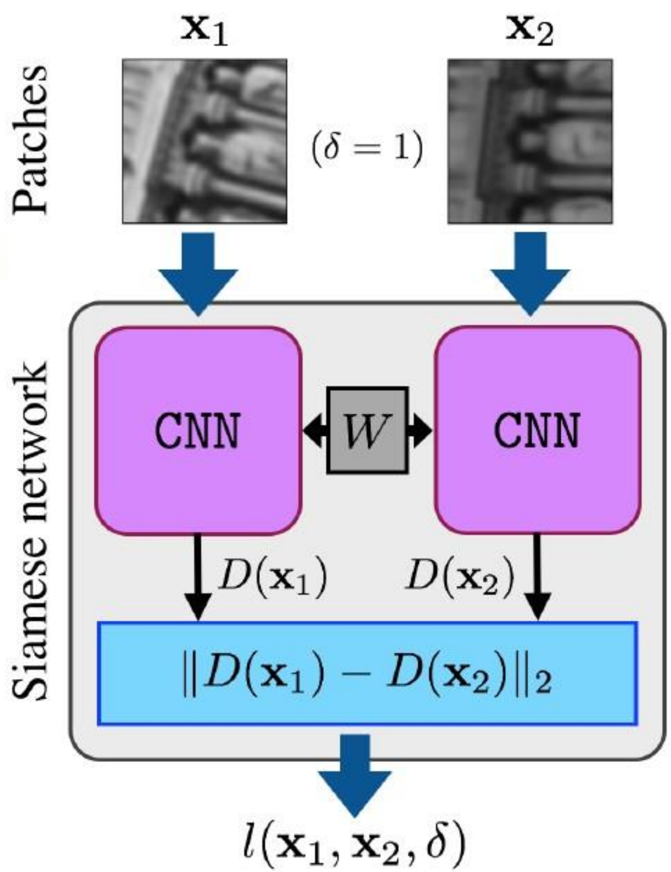

Fracking Deep Convolutional Image Descriptors, http://arxiv.org/abs/1412.6537
Convolutional Neural Networks learn compact local image descriptors, http://arxiv.org/abs/1304.7948

Siamese Network

Stereo

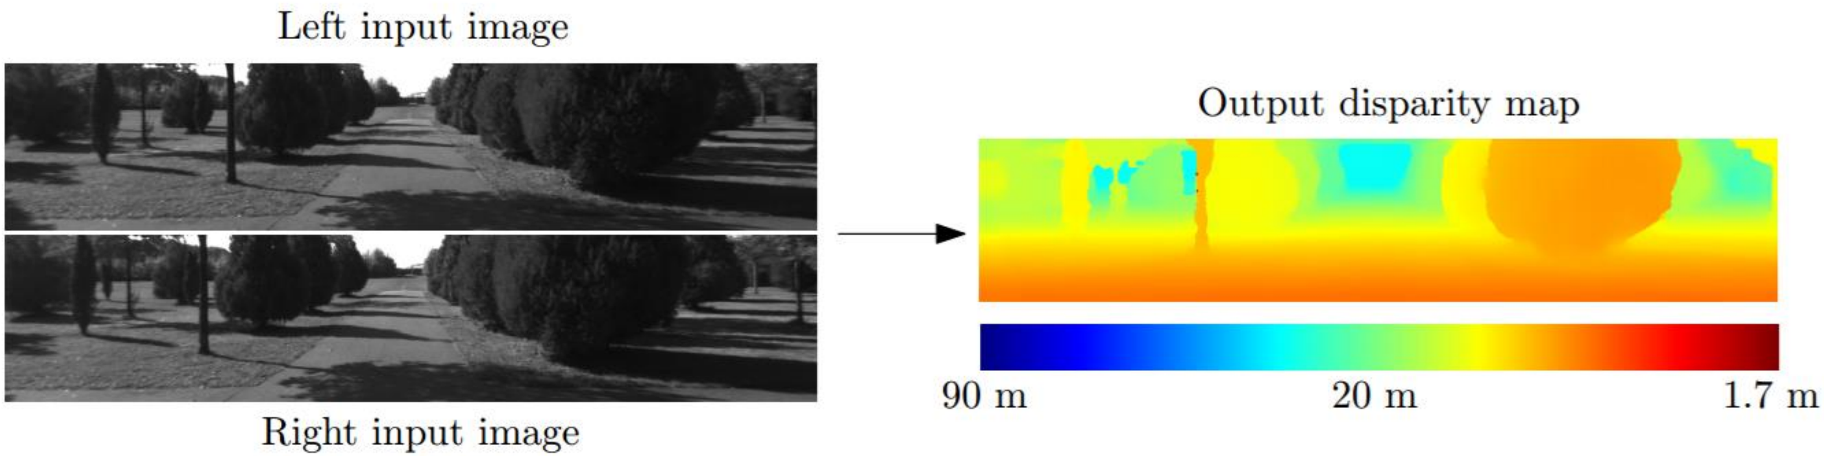

Zbontar, Jure, and Yann LeCun. "Stereo matching by training a convolutional neural network to compare image patches." *Journal of Machine Learning Research* (2016)

Zbontar, Jure, and Yann LeCun. "Computing the stereo matching cost with a convolutional neural network." *CVPR* 2015.

Siamese Network
Stereo

- Correspondances on image locations(Matching)
  - *Good feature*
- Refinement in practice
  - Smoothing

Zbontar, Jure, and Yann LeCun. "Stereo matching by training a convolutional neural network to compare image patches." *Journal of Machine Learning Research* (2016)
Zbontar, Jure, and Yann LeCun. "Computing the stereo matching cost with a convolutional neural network." *CVPR* 2015.

Siamese Network

Stereo

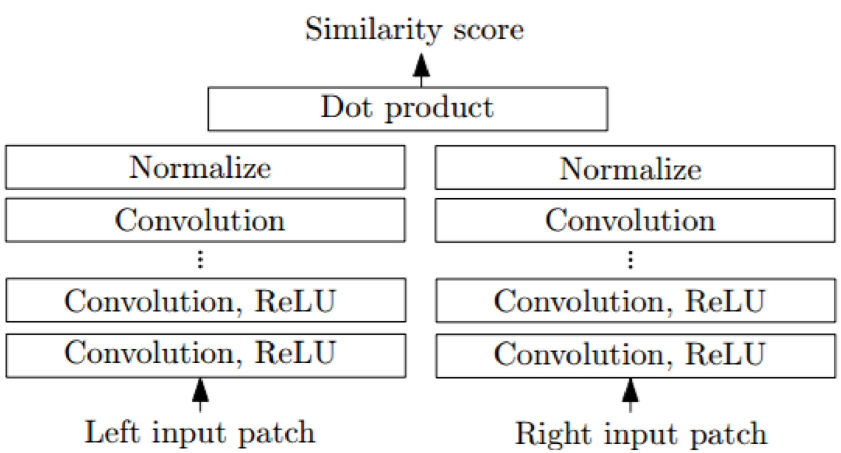

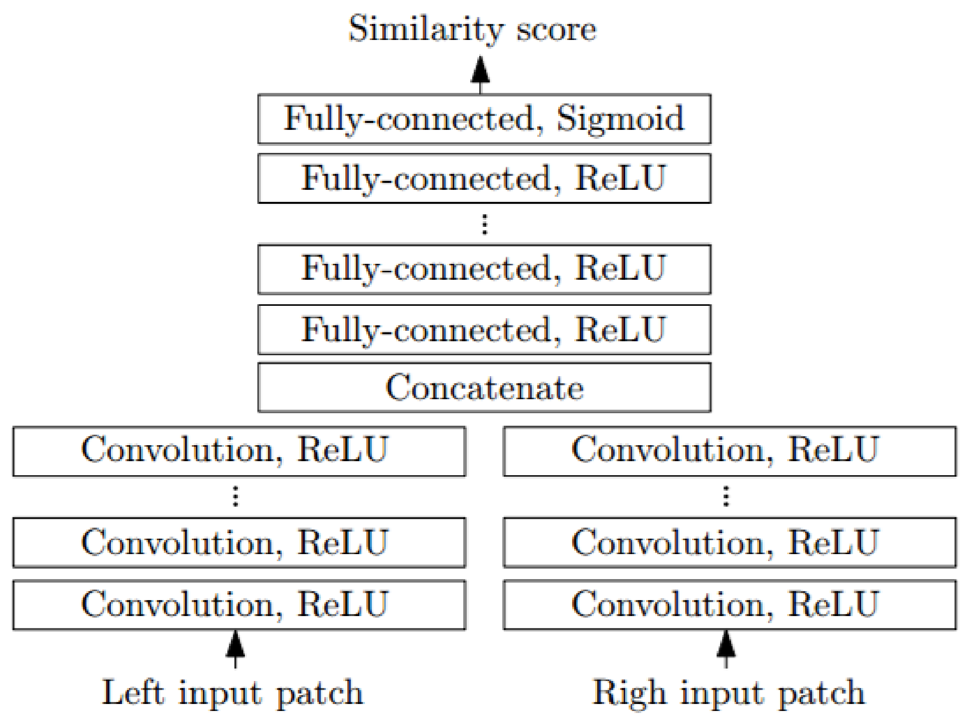

Zbontar, Jure, and Yann LeCun. "Stereo matching by training a convolutional neural network to compare image patches." *Journal of Machine Learning Research* (2016)
Zbontar, Jure, and Yann LeCun. "Computing the stereo matching cost with a convolutional neural network." *CVPR* 2015.

Face recognition

1. Face identification

Who?

A

B

C

...

...

Multiclass classification

2. Face verification

Same?

Same person or not.

Binary Result

http://vis-www.cs.umass.edu/lfw/person/Janica_Kostelic.html

Siamese Network

Application in face verification(I)

Let X_1 and X_2 be a pair of images shown to our learning machine. Let Y be a binary label of the pair, $Y = 0$ if the images X_1 and X_2 belong to the same person (a "genuine pair") and $Y = 1$ otherwise (an "impostor pair").

We assume that the loss function depends on the input and the parameters only indirectly through the energy. Our loss function is of the form:

$$\mathcal{L}(W) = \sum_{i=1}^{P} L(W, (Y, X_1, X_2)^i)$$

$$L(W, (Y, X_1, X_2)^i) = (1 - Y) L_G \left(E_W(X_1, X_2)^i\right) + Y L_I \left(E_W(X_1, X_2)^i\right)$$

$$= (1 - Y)\frac{2}{Q}(E_W)^2 + (Y) 2Q\, e^{-\frac{2.77}{Q} E_W}$$

$$E_W = \|G_W(X_1) - G_W(X_2)\|$$

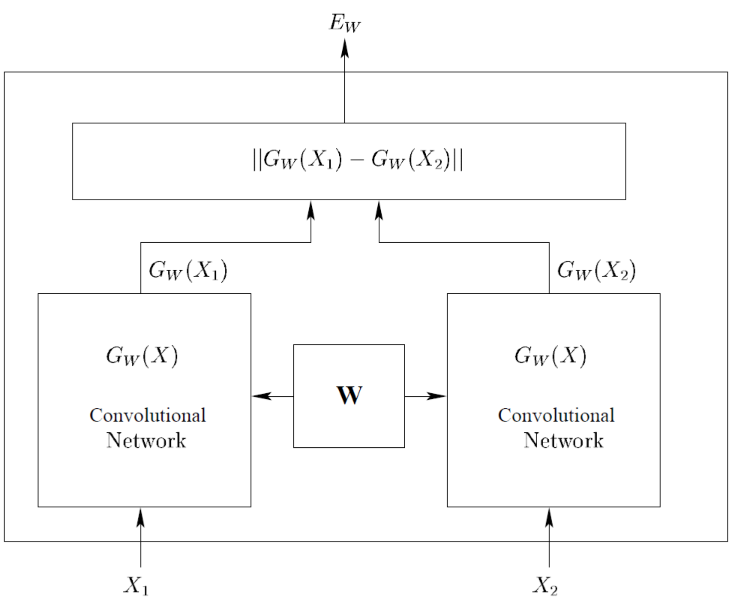

Chopra S, Hadsell R, LeCun Y. Learning a similarity metric discriminatively, with application to face verification, CVPR 2005

Siamese Network

Application in face verification(II) LFW:90.68%

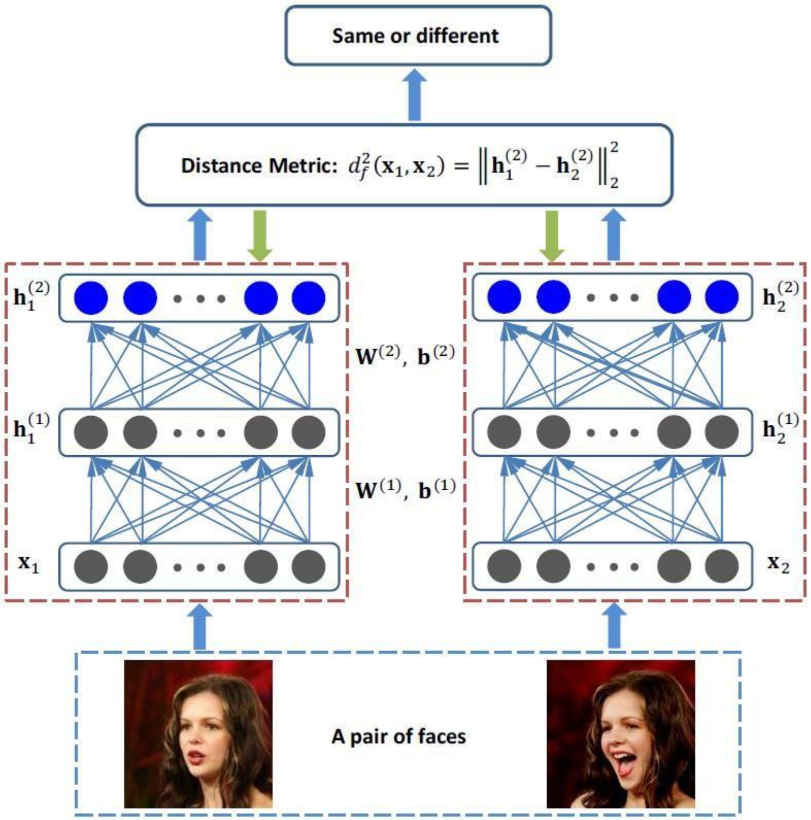

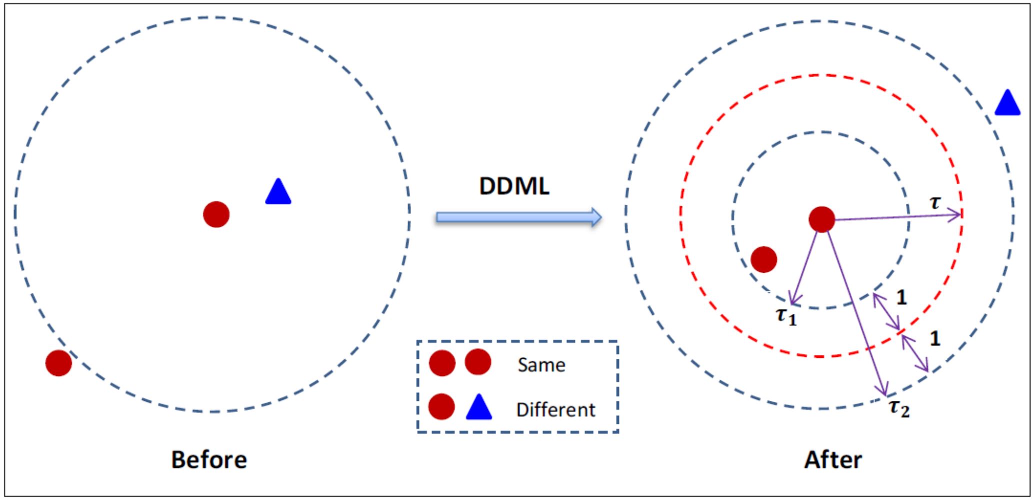

Intuitive illustration of the proposed DDML method

Junlin Hu, etc. Discriminative Deep Metric Learning for Face Verification in theWild, CVPR 2014

Siamese Network

Application in face verification(II)

$$d_f^2(x_i, x_j) < \tau - 1, l_{ij} = 1$$
$$d_f^2(x_i, x_j) > \tau + 1, l_{ij} = -1$$

$$\ell_{ij}\left(\tau - d_f^2(\mathbf{x}_i, \mathbf{x}_j)\right) > 1$$

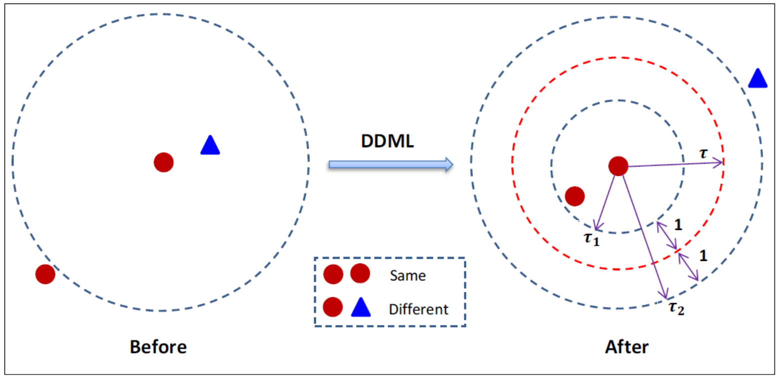

Intuitive illustration of the proposed DDML method

DDML as the following optimization problem:

$$\begin{aligned}\arg\min_f J &= J_1 + J_2 \\ &= \frac{1}{2}\sum_{i,j} g\Big(1 - \ell_{ij}\big(\tau - d_f^2(\mathbf{x}_i, \mathbf{x}_j)\big)\Big) \\ &+ \frac{\lambda}{2}\sum_{m=1}^{M}\Big(\|\mathbf{W}^{(m)}\|_F^2 + \|\mathbf{b}^{(m)}\|_2^2\Big)\end{aligned}$$

where $g(z) = \frac{1}{\beta}\log\big(1 + \exp(\beta z)\big)$ is the generalized logistic loss function [25], which is a smoothed approximation of the hinge loss function $[z]_+ = \max(z, 0)$

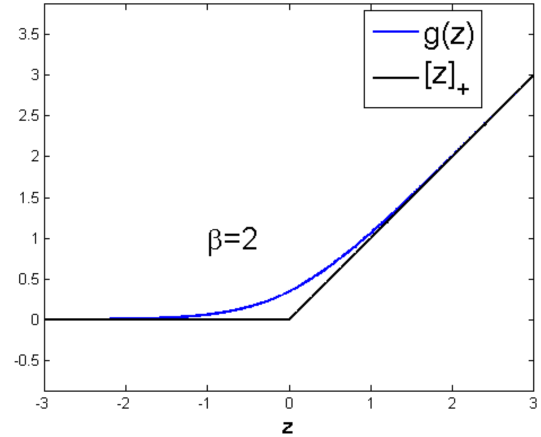

Junlin Hu, etc. Discriminative Deep Metric Learning for Face Verification in theWild, CVPR 2014

Classification Network

Application in face verification(IV)

LFW:97.35%

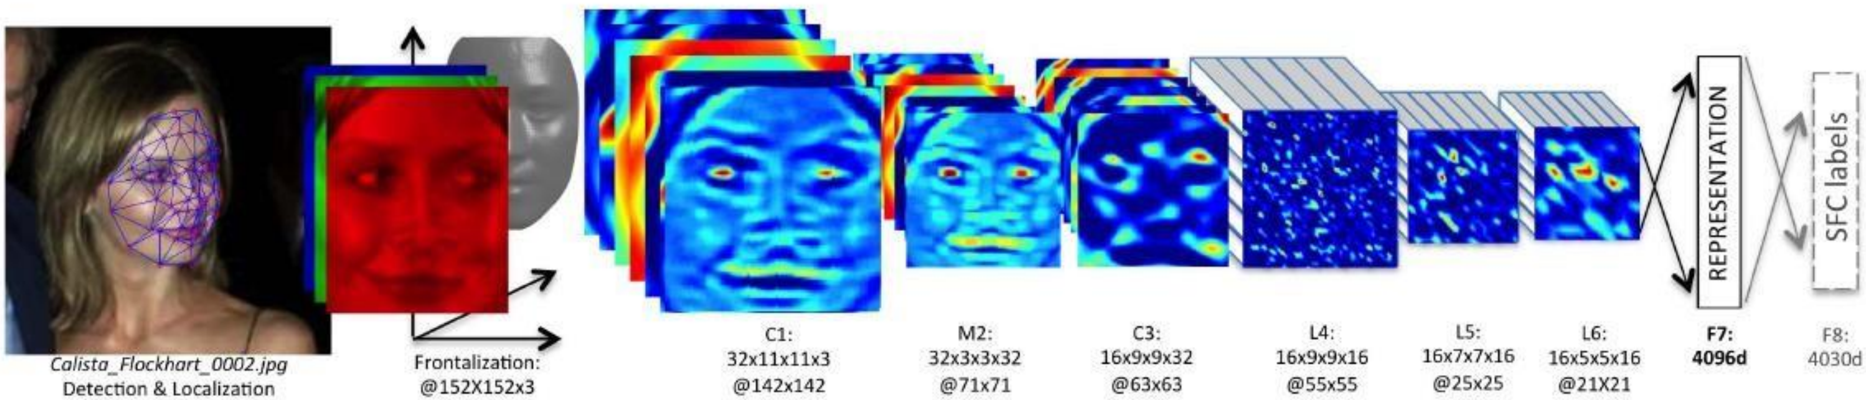

Verification Metric:

1)Cosine similarity

2) Weighted χ2 distance $\chi^2(f_1, f_2) = \sum_i w_i (f_1[i] - f_2[i])^2 / (f_1[i] + f_2[i])$

3) Siamese network $d(f_1, f_2) = \sum_i \alpha_i |f_1[i] - f_2[i]|$

Yaniv Taigman, etc. DeepFace: Closing the Gap to Human-Level Performance in Face Verification, CVPR 2014

Classification Network

Application in face verification(Ⅲ)

LFW:97.45%

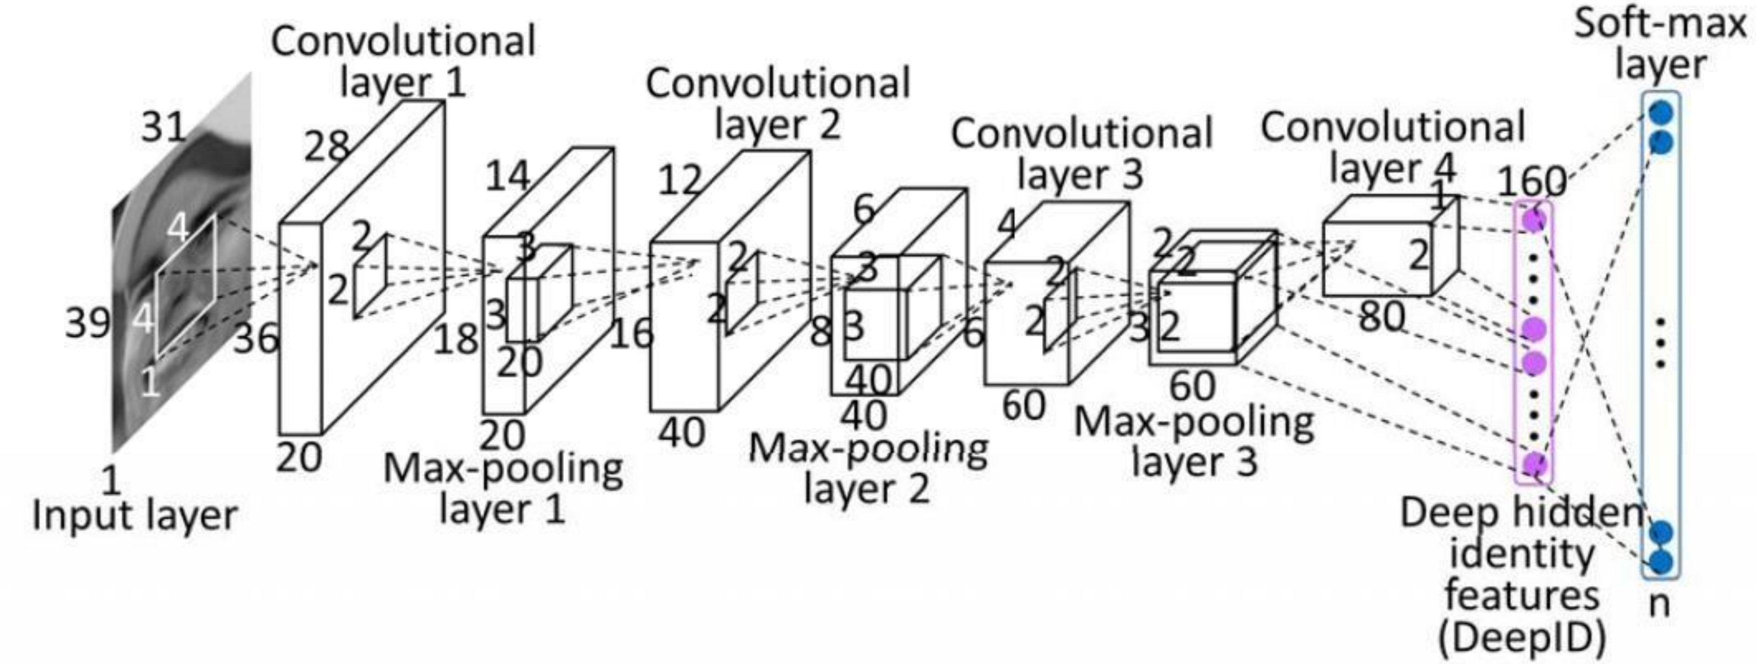

Face Verification: Joint Bayesian

Yi Sun, etc. Deep Learning Face Representation from Predicting 10,000 Classes, CVPR 2014

Classification & Siamese Network

Application in face verification(V)

LFW:99.15%

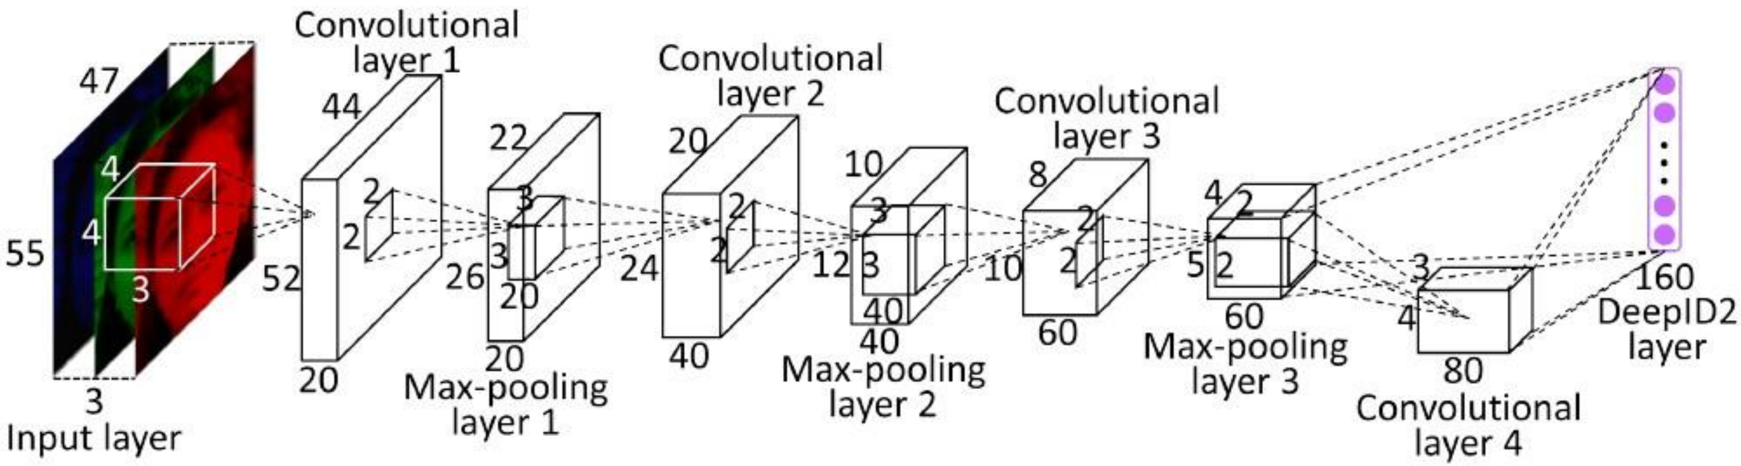

Figure 1: The ConvNet structure for DeepID2 extraction.

Deep Learning Face Representation by Joint Identification-Verification

Yi Sun, etc. Deep Learning Face Representation by Joint Identification-Verification. NIPS 2014

Classification & Siamese Network

Application in face verification(V)

1. identification loss(cross-entropy)

$$\mathrm{Ident}(f, t, \theta_{id}) = -\sum_{i=1}^{n} -p_i \log \hat{p}_i = -\log \hat{p}_t$$

2. verification loss (contrastive)

$$\mathrm{Verif}(f_i, f_j, y_{ij}, \theta_{ve}) = \begin{cases} \frac{1}{2} \|f_i - f_j\|_2^2 & \text{if } y_{ij} = 1 \\ \frac{1}{2} \max\left(0, m - \|f_i - f_j\|_2\right)^2 & \text{if } y_{ij} = -1 \end{cases}$$

3. verification loss (cosine)

$$\mathrm{Verif}(f_i, f_j, y_{ij}, \theta_{ve}) = \frac{1}{2}\left(y_{ij} - \sigma(wd + b)\right)^2$$

where $d = \frac{f_i \cdot f_j}{\|f_i\|_2 \|f_j\|_2}$ is the cosine similarity

Yi Sun, etc. Deep Learning Face Representation by Joint Identification-Verification. NIPS 2014

Triplet Network

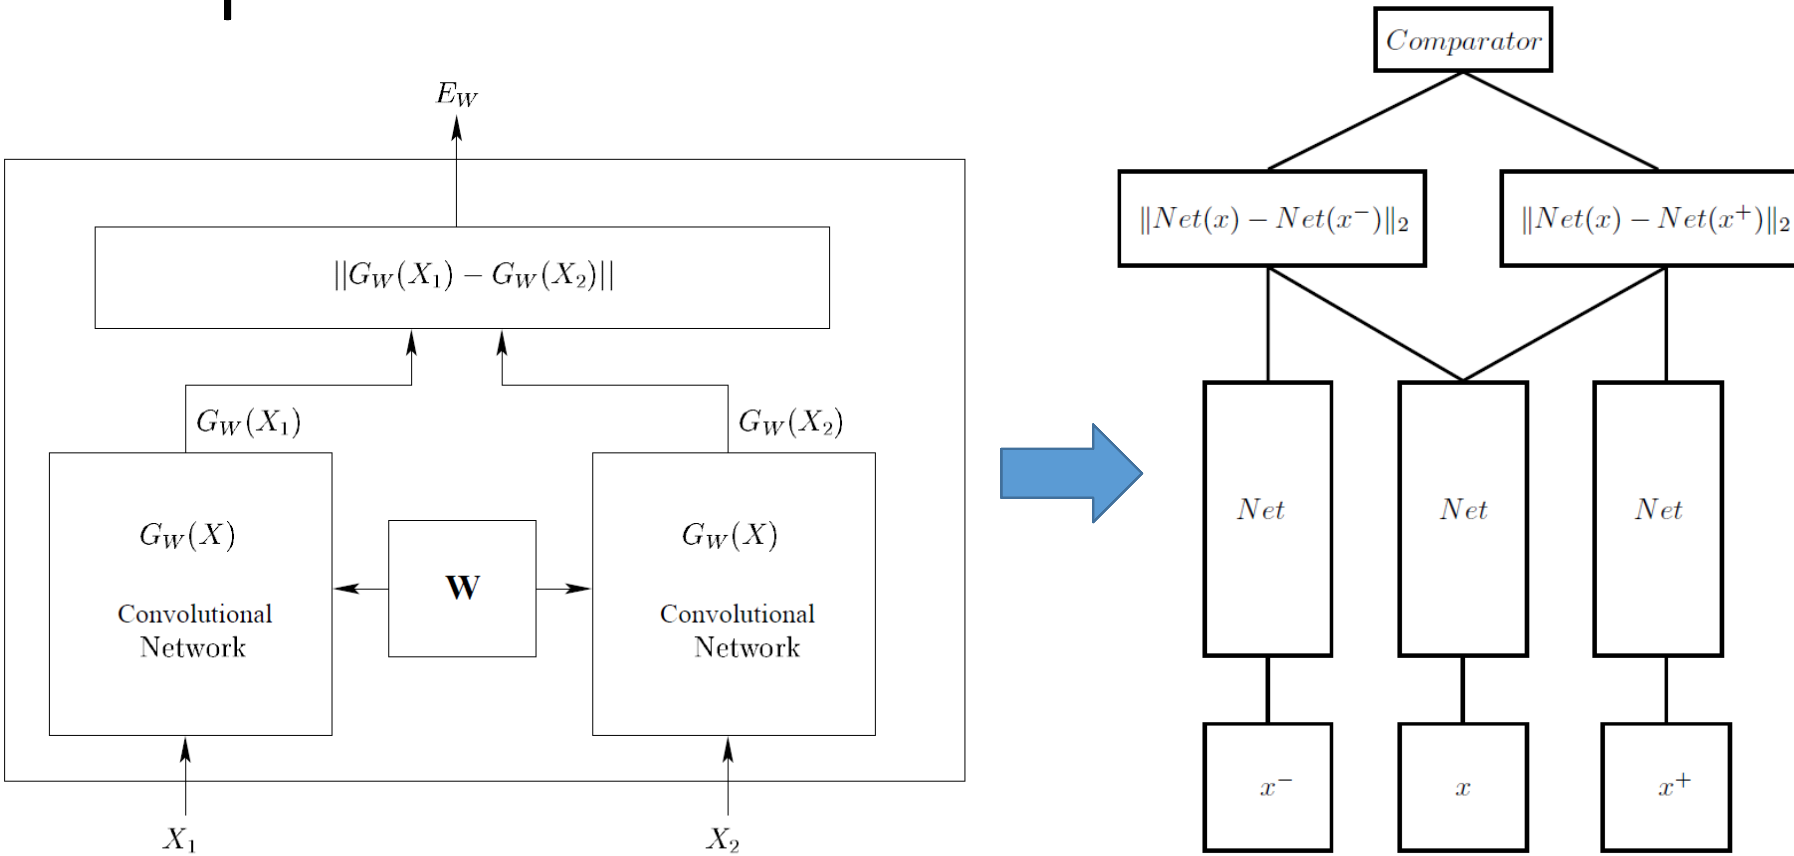

From Siamese to Triplet Network

Elad Hoffer, etc. DEEP METRIC LEARNING USING TRIPLET NETWORK http://arxiv.org/abs/1412.6622

Triplet Network

Application in Image ranking

Sample images from the triplet dataset

Jiang Wang, etc. Learning Fine-grained Image Similarity with Deep Ranking. CVPR 2014

Triplet Network

Application in Image ranking

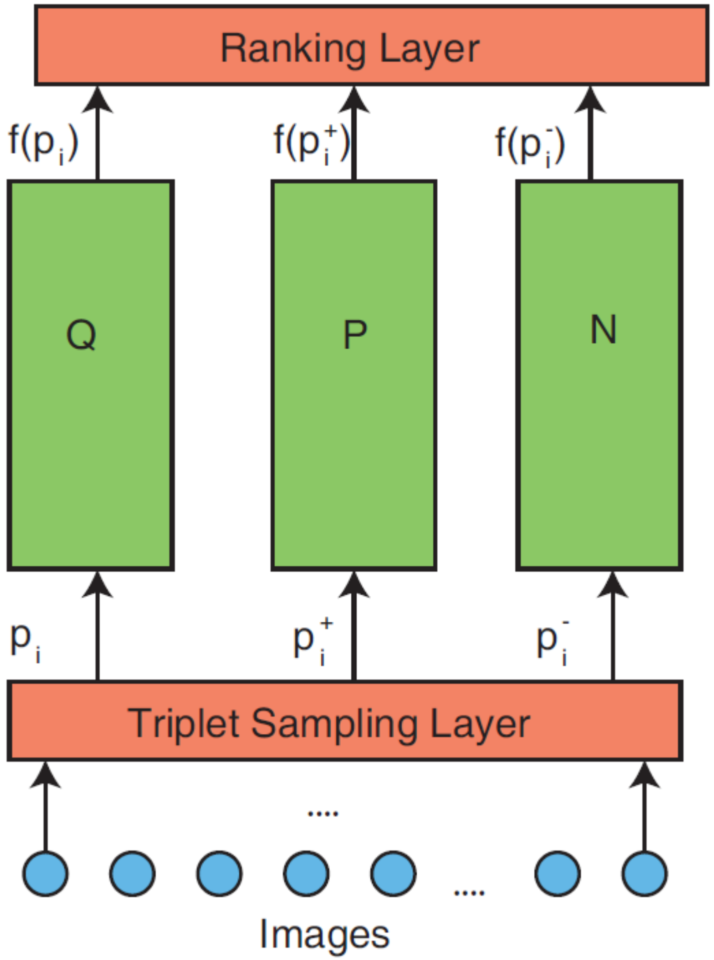

Jiang Wang, etc. Learning Fine-grained Image Similarity with Deep Ranking. CVPR 2014

Triplet Network

Application in Image ranking

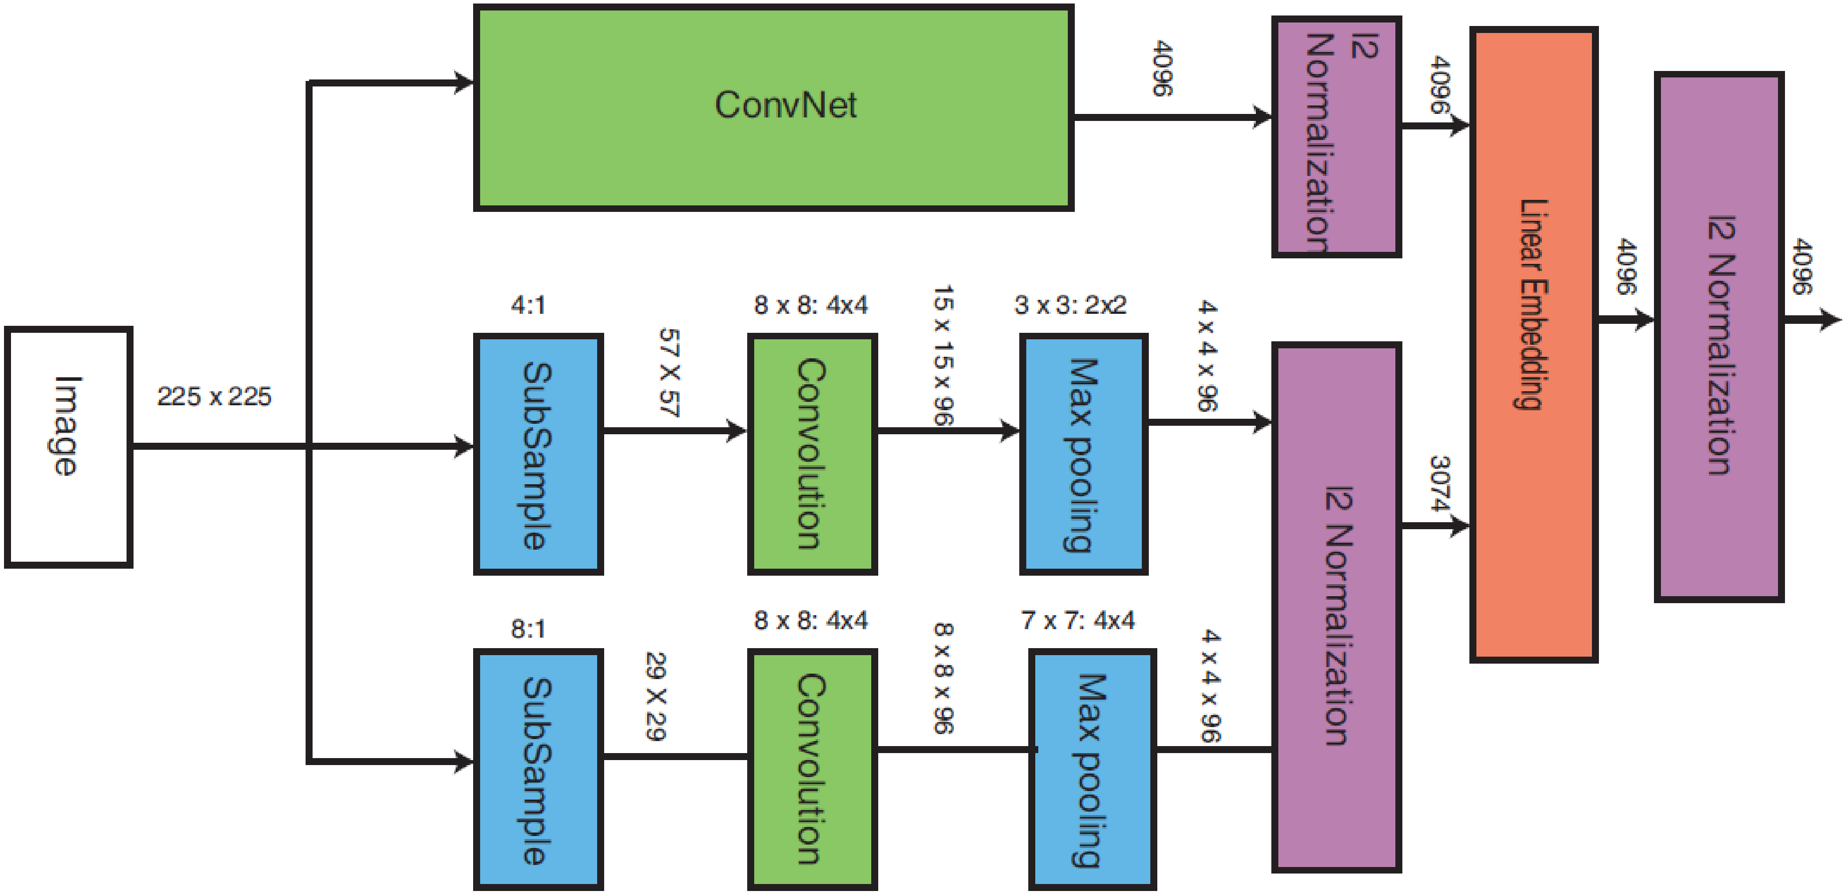

Jiang Wang, etc. Learning Fine-grained Image Similarity with Deep Ranking. CVPR 2014

Triplet Network
Application in Image ranking

Distance

$$D(f(P), f(Q)) = \|f(P) - f(Q)\|_2^2$$

$$D(f(p_i), f(p_i^+)) < D(f(p_i), f(p_i^-)),$$
$$\forall p_i, p_i^+, p_i^- \text{ such that } r(p_i, p_i^+) > r(p_i, p_i^-)$$

Hinge Loss

$$l(p_i, p_i^+, p_i^-) =$$
$$\max\{0, g + D(f(p_i), f(p_i^+)) - D(f(p_i), f(p_i^-))\}$$

Jiang Wang, etc. Learning Fine-grained Image Similarity with Deep Ranking. CVPR 2014

Triplet Network
Application in deep metric learning

$$TripletNet(x, x^-, x^+) = \begin{bmatrix} \|Net(x) - Net(x^-)\|_2 \\ \|Net(x) - Net(x^+)\|_2 \end{bmatrix} \in \mathbb{R}^2_+$$

SoftMax function is applied on both outputs

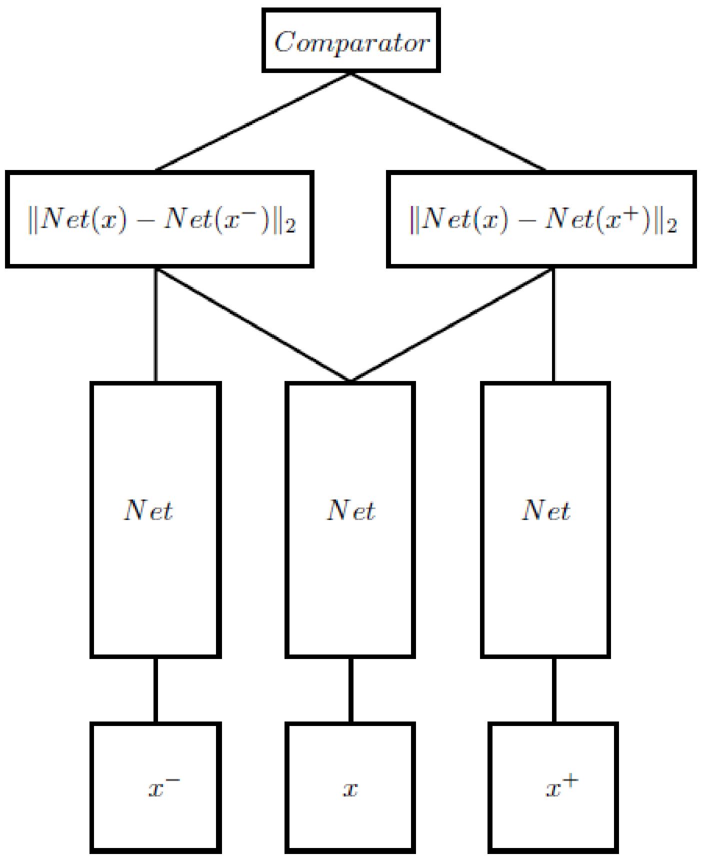

Elad Hoffer, etc. DEEP METRIC LEARNING USING TRIPLET NETWORK. *http://arxiv.org/abs/1412.6622*

Triplet Network

Application in deep metric learning

2D VISUALIZATION OF FEATURES of CIFAR10

Elad Hoffer, etc. DEEP METRIC LEARNING USING TRIPLET NETWORK. [*http://arxiv.org/abs/1412.6622*](http://arxiv.org/abs/1412.6622)

Conclusion

- The loss function in Siamese Network is very important.
- Mixed Network Architecture can improve the performance.
- Caffe implementation of Siamese Network: http://caffe.berkeleyvision.org/gathered/examples/siamese.html

Thank you!